Range-Based Set Synchronization

Technische Universität Berlin

Aljoscha Meyer

September 28, 2021

Range-Based Set Synchronization

Abstract

Set reconciliation is the problem of computing the union between two sets stored on different machines. Set mirroring is the problem of updating one of these sets to be equal to the other. Optimal communication complexity requires transmitting a number of bits proportional to $n_\triangle$, the number of items occurring in exactly one of the two sets. While there are sophisticated probabilistic set reconcilliation protocols achieving this optimum, they require computation time proportional to the size of the sets and computation space proportional to $n_\triangle$. Furthermore, they cannot be modified to solve the set mirroring problem.

Range-based synchronization is a simpler approach based on comparing hashes of ranges over the set and recursing if the two nodes arrive at different hashes, solving both the set reconciliation and the set mirroring problem. The number of bits transmitted is worse than the optimum by a factor of the logarithm of the size of the sets, but in exchange the computation time is only proportional to $n_\triangle$ and the computation space is constant.

Contents

1 Introduction **5**
1.1 Motivating Examples . . . 5
1.2 Efficiency Criteria . . . 6
1.3 Range-Based Synchronization . . . 7
1.4 Contributions and Outline . . . 8

2 Computing Fingerprints **9**
2.1 Initial Considerations . . . 9
2.2 Incremental Computations . . . 11
2.2.1 Subsets . . . 13
2.2.2 Bulk Fingerprint Computation . . . 17
2.2.3 Differences to Homomorphic Hashing . . . 19
2.3 Monoidal Fingerprints . . . 21
2.4 Cryptographically Secure Fingerprints . . . 22
2.4.1 General Considerations . . . 23
2.4.2 Specific Monoids . . . 25
2.5 Pseudorandom Data Structures . . . 27
2.5.1 Pseudorandom Treaps . . . 28
2.5.2 Pseudorandom Skip Lists . . . 31

3 Set Reconciliation **35**
3.1 Recursive Set Reconciliation . . . 35
3.1.1 Observations . . . 36
3.2 Proof of Correctness . . . 37
3.3 Complexity Analysis . . . 38
3.3.1 Preliminary Observations . . . 38
3.3.2 Communication Rounds . . . 38
3.3.3 Communication Complexity . . . 38
3.3.4 Impact of Randomized Range Boundaries . . . 40
3.3.5 Computational Complexity . . . 40
3.3.6 Discussion . . . 41
3.4 Smaller Optimizations . . . 42
3.4.1 Non-Uniform Partitions . . . 42
3.4.2 Subset Checks . . . 42
3.4.3 Efficiently Encoding Ranges . . . 43
3.4.4 Utilizing Range Boundaries . . . 43
3.4.5 Multi-Level Fingerprints . . . 44
3.5 Reconciling Hash Graphs . . . 44

4 Protocol Variants **46**
4.1 Higher-Dimensional Ranges 46
4.2 Map Reconciliation 47
4.3 Set and Map Mirroring 47
4.4 Authenticated Data Structures 48

5 Related Work **49**
5.1 Fingerprinting Sets 49
5.2 Set Reconcilliation 50
5.2.1 Characteristic Polynomial Interpolation 50
5.2.2 Invertible Bloom Lookup Tables 51
5.2.3 Other Bloom Filter Approaches 52
5.2.4 Partition Reconciliation 52

6 Conclusion **54**
6.1 Future Work 54

Chapter 1

Introduction

One of the problems that needs to be solved when designing a distributed system is how to efficiently synchronize data between nodes. Two nodes may each hold a particular set of data, and may then wish to exchange the ideally minimum amount of information until they both reach the same state. Typical ways in which this can happen are one node taking on the state of the other, or both nodes ending up with the union of information available between the two of them. We study the synchronization of sets and key-value mappings by exchanging fingerprints for recursively smaller sub-ranges of the item/key domain. This protocol is stateless, it requires no information about the local data structure to synchronize.

1.1 Motivating Examples

Distributed version control systems can be seen as an example where nodes symmetrically exchange information to update each other's state: users independently create new objects which describe changes to a directory, and when connecting to each other, they both fetch all new updates from the other node in order to obtain a (more) complete version history. Regarded more abstractly, the two nodes compute the union of the sets of objects they store. A version control system might attempt to leverage structured information about those objects, such as a happened-before relation, but this does not lead to good worst-case guarantees. The set reconciliation protocol we give guarantees the exchange to only take a number of rounds logarithmic in the number of objects.

A different example are peer-to-peer publish-subscribe systems such as Secure Scuttlebutt [TLMT19]. A node in the system can subscribe to any number of topics, and nodes continuously synchronize all topics of interest that they share with the other nodes they encounter on a randomized overlay network. Scuttlebutt achieves efficient synchronization by enforcing a linear happened-before relation between messages published to the same topic, i.e. each message is assigned a unique sequence number that is one greater than the sequence number of the previous message in the topic. When two nodes share interest in a topic, they exchange the greatest sequence number they each have for this topic, whichever node sent the greater one then knows which messages the other is missing.

The price to pay for this efficient protocol is that concurrent publishing of new messages to the same topic is forbidden, since that would lead to multiple messages with the same sequence number, breaking correctness of the synchronization proce-

dure. An unordered pubsub mechanism based on set reconciliation would be able to support concurrent publishing, the other design aspects such as the overlay network could be left unchanged.

An example from a less decentralized setting are incremental software updates. A server might host a new version of an operating system, and users running an old version want to efficiently download the changes. An almost identical problem is that of efficiently creating a backup of a file system to a server already holding an older backup. Both of these examples can abstractly be regarded as updating a map from file paths to file contents. Our protocol for mirroring maps could be used for determining which files need to be updated. The protocol allows synchronizing the actual files via an arbitrary nested synchronization protocol, e.g. rsync [TM+96].

1.2 Efficiency Criteria

There are a variety of criteria by which to evaluate a synchronization protocol. We exemplify them by the trivial synchronization protocol, which consists of both nodes immediately sending all their data to the other node.

Let n_0 be the number of items held by node $\mathcal{X}_0$, and n_1 the number of items held by node $\mathcal{X}_1$. Let $n_\triangle$ be the number of items held by exactly one of the two nodes, and let $n := n_0 + n_1$.

The most obvious efficiency criteria are the *total bandwidth* ($\mathcal{O}(n)$ items are transmitted for the trivial protocol) and the number of *communication rounds* ($\mathcal{O}(1)$ for the trivial protocol). Suppose one of the nodes holds no items, then the other node must transmit all of its items, so the total bandwidth is always in $\mathcal{O}(n)$. A more useful point of view is that of stating the communication complexity in terms of the minimum number of items which have to be transmitted, i.e. $n_\triangle$. This is where the unsuitability of the trivial protocol becomes apparent, as it transmits $\mathcal{O}(n)$ items even if $n_\triangle$ is small.

Efficiently using the network is not everything, the computational *time complexity per round-trip* must be feasible so that computers can actually run the protocol. It is lower-bounded by the amount of items sent in a given round, since each of of them has to be transmitted. The trivial protocol achieves this lower bound.

Similarly, the *space complexity per round-trip* plays a relevant role, since computers have only a limited amount of memory. In particular, if an adversarial node can make a node run out of memory, the protocol can only be run in trusted environments. Even then, when non-malicious nodes of vastly differing computational capabilities interact (e.g. a microcontroller connecting to a server farm), "accidental denial-of-service attacks" can easily occur. Since the trivial protocol does not perform any actual computation, its space complexity per round-trip is in $\mathcal{O}(1)$.

The *space complexity per session* measures the amount of state nodes need to store across an entire synchronization session, in particular while idly waiting for a response.

In addition to the space required for per-round-trip computations and per session, an implementation of a protocol might need to store auxiliary information that is kept in sync with the data to be synchronized, in order to achieve sufficiently efficient time and space complexity per round-trip. Of interest is not only the *space for the auxiliary information*, but also its *update complexity* for keeping it synchronized with the underlying data.

A protocol might be asymmetric, with different resource usage for different nodes. If there is client and a clear server role, traditionally protocol designs aim to keep the resource usage of the server as low as possible, motivated by the assumption that many clients might concurrently connect to a single server, but a single client rarely connects to a prohibitive amount of servers at the same time.

A final, "soft" criterium is that of simplicity. While ultimately time and space complexities should guide adoption decisions, complicated designs are often a good indicator that the protocol will never see any deployment. Our designs require merely comparisons of (sums of) hashes, and the auxiliary data structure that enables efficient implementation is a simple balanced tree.

The above criteria are tailored to stateless protocols. A different class of protocols could maintain some state in addition to the data structure in question, for example a history of how it was created, or a summary of the outcomes of past synchronization sessions. Such protocols might achieve better efficiency guarantees, but they are less flexible, or might require space proportional to the number of synchronization peers. This work focuses completely on stateless protocols.

1.3 Range-Based Synchronization

The use-cases given above - distributed version control, unordered publish-subscribe, and update distribution - all involve frequent synchronization with incrementally changing data sets. For these situations, protocols in which each node has a computational overhead proportional to the size of its data structure are not appropriate. Range-based synchronization is an approach that incurs some storage overhead in form of an auxiliary data structure, but in exchange achieves computational cost per synchronization session of $\mathcal{O}(\min(n_\triangle \cdot \log(n), n))$.

We briefly sketch the approach in the form of a set reconciliation protocol (i.e., a protocol for computing the union of two sets on different machines) that exemplifies the core ideas. The protocol leverages the fact that large sets can be partitioned into a number of smaller subsets. The protocol assumes that the sets contain elements from a universe on which there is a total order based that can be used to define ranges, and that nodes can efficiently compute fingerprints for any subrange of the universe.

Suppose, for example, two nodes $\mathcal{A}$, $\mathcal{B}$ each hold a set of natural numbers. They can reconcile all numbers within a range as follows: $\mathcal{A}$ computes a fingerprint over all the numbers it holds within the range and then sends this fingerprint to $\mathcal{B}$, together with the range boundaries. $\mathcal{B}$ then computes the fingerprint over all numbers it holds within that same range. There are three possible cases:

- $\mathcal{B}$ computed the same fingerprint it received, then the range is considered to be fully reconciled and the protocol terminates. This makes the protocol probabilistic, it is only reliable if the probability of two different sets having the same fingerprint is negligible.

- $\mathcal{B}$ has no numbers within the range, $\mathcal{B}$ then notifies $\mathcal{A}$, $\mathcal{A}$ transmits all its numbers from the range, and the range has been fully reconciled.

- Otherwise, $\mathcal{B}$ splits the range into two sub-ranges, such that $\mathcal{B}$ has a roughly equal number of numbers within each range. $\mathcal{B}$ then initiates reconciliation

for both of these ranges, the roles of $\mathcal{A}$ and $\mathcal{B}$ reverse.

Crucially, in the last case, the two recursive protocol invocations can be performed in parallel. The number of parallel sessions increases exponentially, so the original range is being reconciled in a number of rounds logarithmic in the greater number of items held by any node within that range.

1.4 Contributions and Outline

Range-based synchronization has been proposed in [CEG+99], but is only briefly sketched there. We are not aware of any dedicated study of the approach as a viable approach to set reconciliation and related synchronization problems. We strive to fill this gap, providing the following main contributions:

- a precise, stand-alone statement of the protocol
- a proof of correctness and complexity analysis
- an improvement of the per-round and overall computational complexity by a logarithmic factor
- characterizations of suitable fingerprint functions
- a discussion of cryptographically secure fingerprints
- generalizations from set reconciliation to broader synchronization problems

The thesis is structured as follows: This introduction has given a broad overview on the studied problem and has sketched the basic range-based approach. Its viability hinges on the efficient computation of fingerprints, which is thoroughly treated in chapter 2. We then give a definition of the set conciliation protocol in chapter 3, and prove its correctness and its complexity guarantees. Chapter 4 generalizes the approach to related problems. Chapter 5 gives an overview of related work and justifies the chosen approach. We conclude in chapter 6.

Chapter 2

Computing Fingerprints

The protocols we will discuss work by recursively computing and comparing fingerprints of sets. This chapter defines and motivates a specific fingerprinting scheme that admits efficient computation with small overhead for the storage and maintenance of auxiliary data structures. Section 2.1 outlines the solution space and theoretic bounds. Section 2.2 characterizes a family of functions that admit efficient incremental computation, and Section 2.3 proposes members of this family that can be used for fingerprinting, and section 2.4 examine security concerns in the face of malicious parties trying to find fingerprint collisions. We examine randomized solutions which compute with high probability fingerprints in logarithmic time in section 2.5.

2.1 Initial Considerations

Our protocols work by recursively testing fingerprints for equality. For our purposes, we can define a fingerprint or hash function as follows:

Definition 1. A *hash function* is a function $\mathrm{h} : U \to D$ with a finite codomain such that for randomly chosen $u \in U$ and $d \in D$ the probability that $\mathrm{h}(u) = d$ is roughly[1] $\frac{1}{|D|}$. $\mathrm{h}(u)$ is called the *hash of* u, *fingerprint of* u or *digest of* u.

Given a universe U of items, a function $\mathrm{enc} : U \to \{0,1\}^*$ for encoding items as binary strings, a linear order $\preceq$ on U, a hash function $\mathrm{h} : \{0,1\}^* \to \{0,1\}^k$ mapping binary strings to binary strings of length k, and some finite $S \subseteq U$, a natural starting point for defining a fingerprint of the set S is to sort the items according to $\preceq$, concatenate the encodings, and hash the resulting string.

While this is straightforward to specify and implement, it does not suffice for our purposes. To allow for efficient set reconciliation, we need to be able to efficiently compute the new fingerprint after a small modification to the set such as insertion or deletion of a single item. Furthermore, we want to be able to efficiently compute the fingerprints of all subsets defined by a range of the original set.

The fingerprint based on concatenating encodings does not allow for efficient incremental reevaluation. When an item is added to S which is less than any

[1] To keep the focus on data structure synchronization rather than being sidetracked by cryptography, we will for the most part keep arguments about probabilities qualitative rather than quantitative.

item previously in S, the hash function needs to be run over the whole string of length $\mathcal{O}(|S|+1)$ again. Furthermore, for any subrange of the set, a full fingerprint computation needs to be performed as well. Precomputing the fingerprints of all subranges requires a prohibitive amount of space. Every subrange corresponds to a substring of the string consisting of all items in S in ascending order, so there are $\frac{|S|\cdot(|S|+1)}{2} + 1 \in \mathcal{O}(n^2)$ many subranges in total.

The go-to approach for efficiently handling small changes to a set of totally ordered items are (balanced) search trees, we briefly state some definitions.

Definition 2. Let U be a set and $\preceq$ a binary relation on U. We call $\preceq$ a *linear order on U* if it satisfies three properties:

anti-symmetry: for all $x, y \in U$: if $x \preceq y$ and $y \preceq x$ then $x = y$

transitivity: for all $x, y, z \in U$: if $x \preceq y$ and $y \preceq z$ then $x \preceq z$

linearity: for all $x, y \in U$: $x \preceq y$ or $y \preceq x$

If $\preceq$ is a linear order, we write $x \prec y$ to denote that $x \preceq y$ and $x \neq y$.

Definition 3. Let U be a set, $\preceq$ a linear order on U, and $V \subseteq U$. Let T be a rooted directed tree with vertex set V.

Let $v \in V$, then T_v denotes the subtree of T with root v.

T is a *binary search tree on V* if for all inner vertices v with left child a and right child b: $a' \prec v$ for all $a' \in T_a$ and $v \prec b'$ for all $b' \in T_b$.

Definition 4. Let $T = (V, E)$ be a binary search tree and $\varepsilon \in \mathbb{R}_{>0}$. We call T *ε-balanced* if $height(T) \leq \lceil \varepsilon \cdot log_2(|V|) \rceil$. Since the precise choice of ε will not matter for our complexity analyses, we will usually simply talk about *balanced* trees.

In the context of fingerprinting, balanced trees often take the form of Merkle trees [Mer89] — binary trees storing items in their leaves, in which each leaf vertex is labeled with the hash of the associated item, and inner vertices are labeled with the hash of the concatenation of the child labels. The root label serves as a fingerprint for the set of items stored in the leaves.

When inserting or removing an item, the number of labels that need updating is proportional to the length of of the path from the root to the item, so in a balanced tree of n items $\mathcal{O}(log(n))$. The problem with this approach however is that fingerprints are not unambiguous: there are different balanced search trees storing the same items set, and different tree arrangements result in different root labels.

Unfortunately it does not suffice to specify a particular balancing scheme, since different insertion orders of the same overall set of items can result in different trees, even when using the same balancing scheme. While this is sufficient for a setting in which only a single machine updates the set and all other machines apply updates in the same order, as assumed e.g. in [NN00], we aim for a less restrictive setting in which the evolution of the local set does not influence synchronization.

An alternative would be to define exactly one valid tree shape for any set of items, but this precludes logarithmic time complexity of updates, as [Sny77] shows that search, insertion and deletion in such trees take at least $\mathcal{O}(\sqrt{n})$ time in the worst case.

We will inspect two options for circumventing this lower bound and to achieve logarithmic complexity: letting the fingerprint function abstract over the tree shape, downgrading it to an implementation detail, which we examine in section 2.2 to section 2.4, or utilizing randomization to define a unique tree layout which allows logarithmic operations with high probability, which we examine in section 2.5.

2.2 Incremental Computations

We now study a family of fingerprinting functions for sets that admit efficient incremental computation based on an auxiliary tree structure, but whose output does not depend on the exact shape of that tree. We first examine a general class of such functions, which reduce a finite set to a single value according to a monoid.

Definition 5. Let M be a set, $\oplus : M \times M \to U$, and $\mathbb{0} \in M$.

We call $(U, \oplus, \mathbb{0})$ a *monoid* if it satisfies two properties:

associativity: for all $x, y, z \in M$: $(x \oplus y) \oplus z = x \oplus (y \oplus z)$

neutral element: for all $x \in M$: $\mathbb{0} \oplus x = x = x \oplus \mathbb{0}$.

Definition 6. Let U be a set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, and $\mathrm{f} : U \to M$.

We *lift* f *to finite sets via* $\mathcal{M}$ to obtain $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}} : \mathcal{P}(U) \rightharpoonup M$ with:

$$\begin{aligned}
&\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\emptyset) := \mathbb{0} \\
&\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S) := \mathrm{f}(\{\min(S)\}) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{U}}(S \setminus \{\min(S)\})
\end{aligned}$$

In other words, if $S = \{s_0, s_1, \ldots, s_{|S|-1}\}$ with $s_0 \prec s_1 \prec \ldots \prec s_{|S|-1}$, then $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S) = \mathrm{f}(s_0) \oplus \mathrm{f}(s_1) \oplus \ldots \oplus \mathrm{f}(s_{|S|-1})$.

Functions of the form $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$ can be incrementally computed by using labeled binary search trees:

Definition 7. Let U be a set, $S \subset U$ a finite set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, $\mathrm{f} : U \to M$, and let T be a binary search tree on S.

We define a *labeling function* $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}} : S \to M$:

$$\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(v) := \begin{cases}
\mathrm{f}(v), & \text{for leaf } v \\
\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_<) \oplus \mathrm{f}(v) & v \text{ internal vertex with left child } c_< \\
& \text{and no right child} \\
\mathrm{f}(v) \oplus \mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_<) & v \text{ internal vertex with right child } c_> \\
& \text{and no left child} \\
\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_<) \oplus \mathrm{f}(v) \oplus \mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_>) & v \text{ internal vertex with left child } c_< \\
& \text{and right child } c_>
\end{cases}$$

See fig. 2.1 for an example.

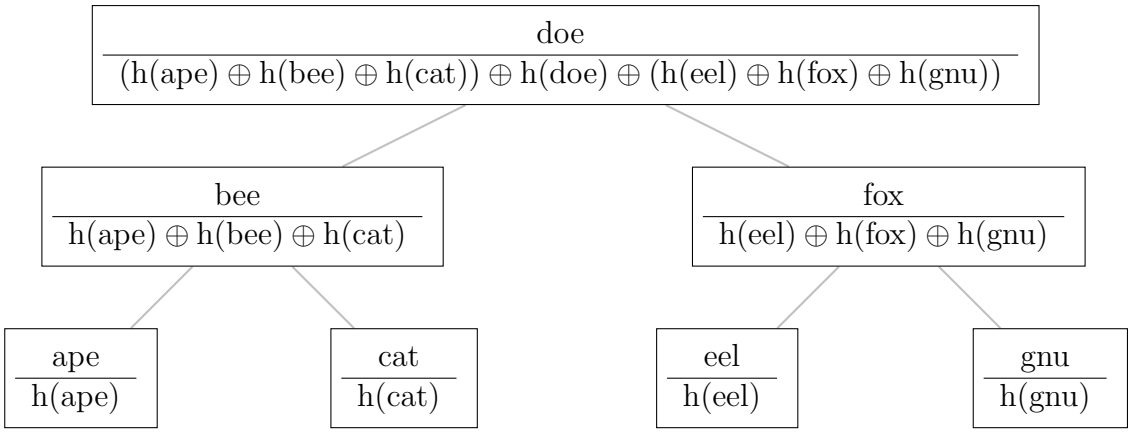

Figure 2.1: A balanced search tree labeled by $\mathrm{label}_{\mathrm{h}}^{(\mathrm{M},\oplus,\mathbb{0})}$. For fingerprinting, h could be a hash function and $\oplus$ the xor operation on fixed-width bitstrings.

Proposition 1. Let U be a set, $S \subset U$ a finite set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, $\mathrm{f} : U \to M$, and let T be a binary search tree on S with root $r \in S$.

Then $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(r) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$.

Proof. By induction on the number of vertices of T.

IB: If r is a leaf, then $|\mathrm{V}(T)| = 1$ and thus $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(r) \overset{\text{def}}{=} \mathrm{f}(r) \overset{\text{def}}{=} \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T)) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$.

IH: Let $c_<$ and $c_>$ be vertices for which $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_<) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T_{c_<}))$ and $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_>) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T_{c_>}))$.

IS: If r is an internal vertex with left child $c_<$ and right child $c_>$, then:

$$\begin{aligned}
\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(r) &\overset{\text{def}}{=} \mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_<) \oplus \mathrm{f}(r) \oplus \mathrm{label}_{\mathrm{f}}^{\mathcal{M}}(c_>) \\
&\overset{\text{IH}}{=} \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T_{c_<})) \oplus \mathrm{f}(r) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T_{c_>})) \\
&\overset{\text{def}}{=} \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\mathrm{V}(T)) \\
&= \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)
\end{aligned}$$

The cases for internal vertices with exactly one child follow analogously. □

This correspondence can be used to incrementally compute $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$: Initially, a labeled search tree storing the items in S is constructed. $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$ is the root label. When an item is inserted or removed, only the labels on the path from the root to the point of modification require recomputation, so only a logarithmic number of operations is performed if a self-balancing tree is used.

Note that the exact shape of the tree determines the grouping of how to apply $\oplus$, but by associativity all groupings yield the same result. All trees storing the same set have the same root label.

If U is small enough that space usage of $\mathcal{O}(|U|)$ is acceptable, an implicit tree representation such as a binary indexed tree (Fenwick tree) [Fen94] can be used. In that case, array positions that correspond to some $u \in U \setminus S$ are simply filled with a dummy value whose hash is defined to be $\mathbb{0}$.

2.2.1 Subsets

In addition to incremental computation of the fingerprint of a given set, the reconciliation protocol also requires the efficient computation of the fingerprints of arbitrary ranges of the given set. We first fix some terminology and notation:

Definition 8. Let $S \subseteq U$, $\preceq$ a linear order over U, and $x, y \in U$.

The *range from x to y in S*, denoted by $[x, y)_S$, is the set $\{s \in S \mid x \preceq s \prec y\}$ if $x \prec y$, and $S \setminus [y, x)_S$ if $y \preceq x$. We call x the *lower boundary* and y the *upper boundary* of the range (even if $y \preceq x$).

Note that the upper boundary is excluded from the range, so in particular $[x, x)_S = S$.

In the remainder of this section, we assume that for all ranges $[x, y)_S$ we have $x \prec y$. If that is not the case, all computations can be performed for the sets $\{s \in S \mid x \preceq s\}$ and $\{s \in S \mid s \prec y\}$ which partition $[x, y)_S$ if $y \preceq x$, and the resulting values can be combined via $\oplus$ to obtain the desired result.

Given a balanced search tree T with root r for a set S that is labeled by $\mathrm{label}_{\mathrm{f}}^{\mathcal{M}}$, we can compute $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S)$ in logarithmic time. Intuitively, one traces paths in T to both x and y, and then the result is the sum over all vertices "in the area between" these paths. For every vertex on the traced paths, the label of the "inner" child vertex summarizes multiple vertices within the area. Summing over all these children yields the value corresponding to the whole inner area. Since the length of the delimiting paths is logarithmic, overall only a logarithmic number of labels needs to be added up.

We now give a precise definition of this algorithm, defining it via Haskell 98 [Jon03] code, not necessarily because it is executable, but mostly to have a typed mathematic notation. For clarity of presentation, only complete binary trees are considered. Arbitrary binary trees can also have inner nodes with exactly one child. These can be handled with almost the same algorithm by acting as if these nodes had a second child which contributes `emptyFingerprint` to any application of `combine`.

We start by specifying the types involved in the computation, being slightly more general than currently necessary so that the algorithm can later be reused for certain non-monoidal fingerprinting schemes:

```
-- U is the type of items, D the codomain of the hash function.
hash :: U -> D

-- For now, the set of fingerprints is identical to the set of hashes.
type Fingerprint = D

-- The fingerprint of the empty set, for now the neutral element.
emptyFp :: Fingerprint
emptyFp = 0

-- Combine three fingerprints into a new fingerprint.
combine :: Fingerprint -> Fingerprint -> Fingerprint -> Fingerprint
combine a b c = a + b + c

-- Convert a hash to a fingerprint, for now the identity function.
```

```haskell
liftFp :: D -> Fingerprint
liftFp d = d

-- A node is either a leaf or an inner vertex.
data Node = Leaf U D | Inner Node U Node D

-- Extract the label of a node.
label :: Node -> D
label Leaf _ fp       = fp
label Inner _ _ _ fp = fp
```

The algorithm proceeds by first finding the vertex v with the smallest distance to the root such that $x \preceq v \prec y$. This might be r itself.

```haskell
-- Find the node within [x, y)_S that is closest to the root,
-- assuming x ≺ y.
findInitial :: Node -> U -> U -> Maybe Node
findInitial (Leaf v _) x y
    | x <= v && v < y = Just v
    | otherwise       = Nothing
findInitial (Inner l v r _) x y
    | v < x           = findInitial r x y
    | v >= y          = findInitial l x y
    | otherwise       = Just v
```

If there is no such v, then $[x, y)_S = \emptyset$ and thus $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S) = \mathbb{0}$. If however there is such a v, then all vertices of T which are not vertices in T_v are either greater than or equal to y if $v \prec x$, or strictly less than x if $x \prec v$; in either case they do not influence $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S)$.

If v is a leaf, $[x, y)_S = \{v\}$ and thus $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S) = \mathrm{f}(v)$. Otherwise, let $c_<$ be the left child of v and let $c_>$ be the right child. Since all vertices in $T_{c_>}$ are greater than v, they are in particular greater than x. Analogously all vertices in $T_{c_<}$ are less than y.

Keeping in mind that $[x, \max(\mathrm{V}(T_{c_<})))_{\mathrm{V}(T_{c_<})}$ simply denotes the set of vertices in $T_{c_<}$ that are strictly greater than x, and analogously $[\min(\mathrm{V}(T_{c_>})), y)_{\mathrm{V}(T_{c_>})}$ denotes the vertices of $T_{c_>}$ that are less than or equal to y, we have:

$$[x, y)_S = [x, \max(\mathrm{V}(T_{c_<})))_{\mathrm{V}(T_{c_<})} \dot\cup \{v\} \dot\cup [\min(\mathrm{V}(T_{c_>})), y)_{\mathrm{V}(T_{c_>})}$$

implying

$$\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, \max(\mathrm{V}(T_{c_<})))_{\mathrm{V}(T_{c_<})}) \oplus \mathrm{f}(v) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([\min(\mathrm{V}(T_{c_>})), y)_{\mathrm{V}(T_{c_>})})$$

```haskell
-- Compute the fingerprint over all items stored in v
-- within the range [x, y)_S, assuming x ≺ y.
rangeFingerprint :: Node -> U -> U -> Fingerprint
rangeFingerprint v x y = case findInitial v x y of
    Nothing                -> emptyFp
```

```
    Just (Leaf _ fp)       -> liftFp 'sfp
    Just (Inner l v' r _) -> combine
                                (accGeq l x) -- liftf^M(u ∈ V(T_l) | x ⪯ u)
                                (liftFp (hash v')) -- f(v')
                                (accLt r y)  -- liftf^M(u ∈ V(T_l) | u ≺ y)
```

$\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, \max(\mathrm{V}(T_{c_<}))]_{\mathrm{V}(T_{c_<})})$ can be computed by traversing from $c_<$ to x, summing over those vertices along the path that are greater or equal to x, as well as the right children of all vertices along the path. Correctness follows from a rather technical but straightforward induction we omit.

```
-- Sum up the fingerprints of all items in the given tree
-- which are greater than or equal to x.
accGeq :: Node -> U -> Fingerprint
accGeq (Leaf v fp) x
    | v < x     = emptyFp
    | otherwise = liftFp fp
accGeq (Inner l v r _) x
    | v < x     = accGeq r x
    | otherwise = combine (accGeq l x) (liftFp (hash v)) (label r)
```

Analogously $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([\min(\mathrm{V}(T_{c_>})), y)_{\mathrm{V}(T_{c_>})})$ can be computed by traversing from $c_>$ to y, summing over those vertices along the path that are strictly less than y, as well as the left children of all vertices along the path.

```
-- Sum up the fingerprints of all items in the given tree
-- which are strictly less than y.
accLt :: Node -> U -> Fingerprint
accLt (Leaf v fp) y
    | v >= y    = emptyFp
    | otherwise = liftFp fp
accLt (Inner l v r _) y
    | v >= y    = accLt l y
    | otherwise = combine (label l) + (liftFp (hash v)) + (accLt r y)
```

A simplified approach to these computations can be taken if M has efficiently computable inverses with respect to $\oplus$, i.e. if $(M, \oplus, \mathbb{0})$ is a group.

Definition 9. Let $(M, \oplus, \mathbb{0})$ be a monoid. We call it a *group* if for all $x \in M$ there exists $y \in M$ such that $x \oplus y = \mathbb{0}$. This y is necessarily unique and denoted by $-x$. For $x, y \in M$ we write $x \ominus y$ as a shorthand for $x \oplus -y$.

Observe that $[x, y)_S = [\min(S), y)_S \setminus [\min(S), x)_S$, and thus also $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([x, y)_S) = -\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([\min(S), x)_S) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}([\min(S), y)_S)$. We can thus simplify `rangeFingerprint`, requiring neither `findInitial` nor `accGeq`:

```
rangeFingerprintGroup :: Node -> U -> U -> Fingerprint
rangeFingerprintGroup v x y = -(accLt v x) + (accLt v y)
```

Complexity

The worst-case running time occurs when the vertex v with the smallest distance to r such that $x \preceq v \prec y$ is r itself, since then both `sumGeq` and `sumLt` perform a traversal to a leaf of maximal length. Since T is balanced, only a logarithmic number of recursive calls is executed. Assuming f can be computed in $\mathcal{O}(1)$, the resulting time complexity is in $\mathcal{O}(\log(|S|))$.

Neither `accLt` nor `accGeq` are tail-recursive, but they can be rewritten into a tail-recursive form via the standard technique of adding an accumulator argument:

```
accGeq :: Node -> U -> Fingerprint
accGeq v x = accGeq' emptyFp v x

accGeq' :: Fingerprint -> Node -> U -> Fingerprint
accGeq' acc (Leaf v fp) x
    | v < x     = acc
    | otherwise = combine emptyFp (liftFp fp) acc
accGeq' acc (Inner l v r _) x
    | v < x     = accGeq' acc r x
    | otherwise = accGeq' (
                      combine (liftFp (hash v)) (label r) acc
                  ) l x

accLt :: Node -> U -> Fingerprint
accLt v y = accLt' emptyFp v y

accLt' :: Fingerprint -> Node -> U -> Fingerprint
accLt' acc (Leaf v fp) y
    | v >= y    = acc
    | otherwise = combine acc (liftFp fp) emptyFp
accLt' acc (Inner l v r _) y
    | v >= y    = accLt' acc l y
    | otherwise = accLt' (
                      combine acc (label l) + (liftFp (hash v))
                  ) r y
```

The space complexity of computing the fingerprint for a range is thus $\mathcal{O}(1)$. Note however that the tail-recursive forms require associativity to compute the same fingerprints as the original versions.

As a closing remark on these computations, we want to point out that associativity not only allows us to abstract over differently shaped binary search trees, but also results in trees of higher degree leading to the same fingerprints. Binary trees simplify the presentation and are theoretically optimal, but on actual hardware traversing pointers is expensive whereas loading sequential memory is cheap, which becomes even more pronounced when the tree resides on secondary storage.

Optimized tree implementations thus have higher branching and store multiple items per vertex, a typical example being B-trees [BM02]. Our fingerprinting scheme admits such optimized implementations, in particular different nodes that want to synchronize data can choose their implementation techniques fully independently and yet maintain compatibility.

2.2.2 Bulk Fingerprint Computation

The synchronization protocols partitions item sets into disjoint ranges. When computing the fingerprints of k such ranges using the previous algorithms, the overall time complexity is in $\mathcal{O}(k \cdot \log(|S|))$. The reason why these algorithms lead to a worse complexityworse than the overall $\mathcal{O}(|S|)$ vertices and edges in the tree, is that there is some amount of redundancy in these independent computations. For example, every single computation touches the root node, roughly half of them touch each of the root's children, etc.

We can bound the complexity of multiple fingerprint computation for disjoint ranges by $\mathcal{O}(|S|)$ by performing a more efficient tree traversal which visits any given vertex at most twice. This is achieved by starting at the lower boundary of the first range, then tracing the (unique) path to the upper boundary of the first range, which is also the lower boundary of the second range, so from here the path to the upper boundary of the second range can be traced, and so on.

Thus we consider a traversal algorithm that traces the shortest path from some vertex x to another vertex y with $x \preceq y$. We express the traversal as a function that takes as arguments some vertex y to search for and a starting vertex x, and that then returns either the next vertex on the shortest path from x to the least vertex z such that $y \preceq z$, or `Nothing` if no such z exists or the traversal has finished, i.e. $x = z$.

Since we use the traversal for the computation of fingerprints over successive ranges, the tree is always traversed from a lesser to a strictly greater node. During a traversal, the function can be called with $y \preceq x$, but only if x lies on the shortest path from some $w \prec y$ to z. The definition we give does not correctly handle inputs for which this does not hold, it cannot perform backward traversals.

If the function is to trace of the shortest path in $\mathcal{O}(1)$ space, the decision which vertex to move to must be based on local information only. A vertex x however does not store enough information to determine whether the target vertex y can be reached by a downward traversal starting at x or whether one first needs to find a common ancestor of x and y. To enable this decision, we store in each vertex a reference to its greatest descendent.

To allow upward traversal in the tree in $\mathcal{O}(1)$ space, we store in each vertex a reference to its parent, if one exists. As usual, we restrict the presentation of the algorithm to complete binary trees.

```haskell
-- Second-to-last member of inner nodes is a pointer to the greatest
-- descendent, last member of all nodes is the parent pointer.
data Node = Leaf U D (Maybe Node) | Inner Node U Node D Node (Maybe Node)

value :: Node -> U
value Leaf v _ _ = v
value Inner _ v _ _ _ _ _ = v

greatestDescendent :: Node -> Node
greatestDescendent v@(Leaf _ _ _) = v
greatestDescendent Inner _ _ _ _ _ g _ = g

-- Return the next vertex on the shortest path
```

```
-- from x to the least vertex z with y ≼ z,
-- assuming there exists w ≺ y such that x lies
-- on the shortest path from w to z.
next :: Node -> U -> Maybe Node
next (Leaf _ _ p) y
    | x < y  = p
    | x >= y = Nothing
next (Inner l x r _ g p) y
    | (value g) < y = p
    | y < x && ((value (greatestDescendent l)) >= y)  = Just l
    | y > x  = Just r
    | otherwise = Nothing -- traversal done
```

Since paths in trees are unique, this traversal touches the same vertices as `rangeFingerprint`, except those visited during the execution of `findInitial`. The fingerprint can thus also be computed during the traversal, by maintaining a single bit of state indicating whether the traversal is in the upward or downward phase:

```
data Phase = Up | Down

isPeak :: Node -> U -> Bool
isPeak src@(Leaf _ _ p) y = (next src y) == p
isPeak src@(Inner _ _ _ _ _ p) y = (next src y) == p

-- Compute the fingerprint for the range [x, y)_S,
-- and also return the last vertex of the traversal.
-- The accumulator must initially be emptyFp, the phase Up,
-- and the starting vertex of the traversal x.
computeFingerprint :: D -> Phase -> Node -> U -> U -> (D, Node)
-- Phase Up correspondends to computing accGeq.
computeFingerprint acc Up src@(Leaf v fp _) x y =
     case (next src y) of
        Nothing
            | v < x = (acc, src)
            | otherwise = (combine acc (liftFp fp) emptyFp, src)
        Just n
            | v < x = computeFingerprint acc Up n x y
            | otherwise = computeFingerprint
                (combine acc (liftFp fp) emptyFp) Up n x y
computeFingerprint acc Up src@(Inner _ v r _ _ _) x y =
    case (next src y) of
        Nothing
            | v < x = (acc, src)
            | otherwise = (combine acc (liftFp (hash v)) emptyFp, src)
        Just n
            | isPeak src y = computeFingerprint
                (combine acc (liftFp (hash v)) emptyFp) Down n x y
            | otherwise && v < x = computeFingerprint acc Up n x y
            | otherwise = computeFingerprint
```

```
                (combine acc (liftFp (hash v)) (label r)) Up n x y
-- Phase Down correspondends to computing accLt.
computeFingerprint acc Down src@(Leaf v fp _) x y =
    | v >= y = (acc, src)
    | otherwise = (combine acc (liftFp fp) emptyFp, src)
computeFingerprint acc Down src@(Inner l v _ _ _ _) x y =
    case (next src y) of
        Nothing
            | v >= y = (acc, src)
            | otherwise = ((combine
                    acc
                    (liftFp (label l))
                    (liftFp (hash v))
                ), src)
        Just n
            | v >= y -> computeFingerprint acc Down n x y
            | otherwise -> computeFingerprint
                (combine
                    acc
                    (liftFp (label l))
                    (liftFp (hash v))
                ) Down n x y
```

By computing the fingerprints of successive ranges via this function, using the output vertex as the starting point for the next traversal, the fingerprints are computed while traversing any edge at most two times (once during a downward traversal, and once during an upward traversal). The overall time complexity for k ranges is thus in $\mathcal{O}(\min(k \cdot \log(|S|), |S|))$. Since the function is tail recursive, the space overhead is in $\mathcal{O}(1)$.

2.2.3 Differences to Homomorphic Hashing

The construction of lifting a function to sets resembles that of multiset homomorphic hash functions, introduced in [CDVD+03], with further instantiations given in [CNQ09] and [MSTA17]. We briefly define multiset homomorphic hashing in order to point out salient differences to our construction.

Definition 10. Let $\mathcal{U}_0 := (U_0, \oplus_0, \mathbb{0}_0)$ and $\mathcal{U}_1 := (U_1, \oplus_1, \mathbb{0}_1)$ be monoids and let $f : U_0 \to U_1$.

We call f a *monoid homomorphism from* $\mathcal{U}_0$ *to* $\mathcal{U}_1$ if it satisfies two properties:

preserves operation: for all $x, y \in U_0$: $\mathrm{f}(x \oplus_0 y) = \mathrm{f}(x) \oplus_1 \mathrm{f}(y)$

preserves neutral element: $f(\mathbb{0}_0) = \mathbb{0}_1$

The second property directly follows from the first one by considering the cases $x := \mathbb{0}_0$ and $y := \mathbb{0}_0$.

Definition 11. Let $\mathcal{S} := (\mathbb{N}^U, \cup, \emptyset)$ be the monoid of multisets over the universe U under union, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, and $\mathrm{f} : \mathbb{N}^U \to M$.

We call f a *multiset homomorphic hash function* if f is a hash function and a monoid homomorphism from $\mathcal{S}$ to $\mathcal{M}$.

A multiset homomorphic hash function allows to incrementally maintain the hash of a multiset under insertions using only $\mathcal{O}(1)$ storage space, by storing the hash of the multiset and combining it with the hash of any inserted item. This can be generalized to support deletions as well: by allowing negative multiplicities for multiset items, multisets form a group under union. If the hash function is a group homomorphism, the hash of the multiset can be updated on deletion by subtracting the hash of the deleted item.

This is more efficient than our approach, which requires $\mathcal{O}(n)$ space for the tree structure. This efficiency comes however at the cost of a loss in flexibility. As multiset union is commutative, so must be the homomorphic monoid/group of fingerprints. Since our protocols require the ability to compute fingerprints over arbitrary ranges and we need the tree of size $\mathcal{O}(n)$ to do so anyways, homomorphic hashing would not yield any space usage benefits, and we thus choose the approach that does not require commutativity.

We additionally want to explicitly work with regular sets, not multisets. The reason why homomorphic hashes operate on multisets is that they admit more natural homomorphisms under union. For example, taking the size of a multiset is a homomorphism into the additive monoid on the natural numbers, whereas this is not true for regular sets, as e.g. $|\{x, y\} \cup \{y, z\}| \neq |\{x, y\}| + |\{y, z\}|$. Set union is however well behaved with respect to taking the size of the sets if it is a disjoint union, i.e. no item occurs in both sets.

For an algebraic-flavored characterization of our fingerprint functions, we thus need to account for items occurring on both sides of the union. This alone still would not change that the resulting concept would necessarily require a commutative monoid, so we must also restrict possible inputs to be sorted. This leads to the following definition:

Definition 12. Let U be a set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, and $\mathrm{f} : \mathcal{P}(U) \rightharpoonup M$ a partial function mapping all finite subsets of U into M.

We call f a *set union somethingsomethingmorphism* if for all finite sets $S_0, S_1 \in \mathcal{P}(U)$ such that $\max(S_0) \prec \min(S_1)$: $\mathrm{f}(S_0 \cup S_1) = \mathrm{f}(S_0) \oplus \mathrm{f}(S_1)$.

This definition captures exactly the functions of form $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$, as shown in the following propositions:

Proposition 2. Let U be a set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, and $\mathrm{f} : U \to M$.

Then $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$ is a set union somethingsomethingmorphism.

Proof. Let $S_0, S_1 \in \mathcal{P}(U)$ be finite sets such that $\max(S_0) \prec \min(S_1)$. Then:

$$\begin{aligned}
\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S_0 \cup S_1) &= \bigoplus_{\substack{s_i \in S_0 \cup S_1, \\ \text{ascending}}} \mathrm{f}(s_i) \\
&= \bigoplus_{\substack{s_i \in S_0, \\ \text{ascending}}} \mathrm{f}(s_i) \oplus \bigoplus_{\substack{s_i \in S_1, \\ \text{ascending}}} \mathrm{f}(s_i) \\
&= \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S_0) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S_1)
\end{aligned}$$

Proposition 3. Let U be a set, $\preceq$ a linear order on U, $\mathcal{M} := (M, \oplus, \mathbb{0})$ a monoid, and $\mathrm{g} : \mathcal{P}(U) \rightharpoonup M$ a set union somethingsomethingmorphism.

Then there exists $\mathrm{f} : U \to M$ such that $\mathrm{g} = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$.

Proof. Define $\mathrm{f} : U \to M$ as $\mathrm{f}(u) := \mathrm{g}(\{u\})$. We show by induction on the size of $S \subseteq U$ that $\mathrm{g}(S) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$.

IB: If $S = \emptyset$, then $\mathrm{g}(S) = \mathbb{0} = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$. Suppose that $\mathrm{g}(\emptyset) \neq \mathbb{0}$, this would contradict the fact that for all $x \in U$ we have $\mathrm{g}(\{x\}) = \mathrm{g}(\{x\}) \oplus \mathrm{g}(\emptyset) = \mathrm{g}(\emptyset) \oplus \mathrm{g}(\{x\})$, which holds because $\{x\} = \{x\} \cup \emptyset = \emptyset \cup \{x\}$ and g is a set union somethingsomethingmorphism.

If $S = \{x\}$, then $\mathrm{g}(S) = \mathrm{f}(x) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)$.

IH: For all sets T with $|T| = n$ it holds that $\mathrm{g}(T) = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(T)$.

IS: Let $S \subseteq U$ with $|S| = n + 1$, then:

$$\begin{aligned}
\mathrm{g}(S) &= \mathrm{g}(\{\min(S)\}) \oplus \mathrm{g}(S \setminus \{\min(S)\}) \\
&\overset{\text{IH}}{=} \mathrm{g}(\{\min(S)\}) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S \setminus \{\min(S)\}) \\
&= \mathrm{f}(\min(S)) \oplus \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S \setminus \{\min(S)\}) \\
&= \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(S)
\end{aligned}$$

As g is only defined over finite inputs, we thus have $\mathrm{g} = \mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$. □

2.3 Monoidal Fingerprints

Now that we have characterized a family of functions that admit efficient recomputation in response to changes to the underlying set as well as efficient computation for ranges within the set, the remaining task is to find such functions which are also suitable fingerprints. This consists of deciding on the monoid of fingerprints, and choosing the mapping from items to monoid elements.

As the fingerprint of a singleton set $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\{u\})$ is equal to $\mathrm{f}(u)$, f must itself already be a hash function. Typical hash functions map values to bit strings of a certain length, i.e. the codomain is $\{0, 1\}^k$ for some $k \in \mathbb{N}$. We will thus consider monoids whose elements can be represented by such bit strings.

A natural choice of the monoid universe is then $[0, 2^k)_{\mathbb{N}}$, some simple monoidical operations on this universe include bitwise xor, addition modulo 2^k, and multiplication modulo 2^k. In the following, addition and multiplication will always be implicitly taken modulo 2^k. Note that xor and addition also admit inverses, so the slightly simplified computation of fingerprints can be used.

Of these three options, multiplication is the least suitable, because multiplying any number by 0 yields 0. Consequently for every set containing an item u with $\mathrm{f}(u) = 0$ the fingerprint of the set would also be 0, which clearly violates the criterion that all possible values for fingerprints occur with equal probability.

Addition and xor however are particularly well-behaved in that regard, as they form finite commutative groups:

Proposition 4. Let $\mathcal{G} := (G, \oplus, \mathbb{0})$ be a finite commutative group, i.e. a group with a finite universe such that for all $x, y \in G$: $x \oplus y = y \oplus x$. Let U be a set and let $f : U \to G$ be a hash function.

Then $\mathrm{lift}_{\mathrm{f}}^{\mathcal{G}}$ is a hash function as well.

Proof. We first show that for randomly chosen $x, y, z \in G$ the probability that $x \oplus y = z$ is $\frac{1}{|G|}$.

For $x, z \in G$ there is $y \in G$ such that $x \oplus y = z$, namely $y := z \ominus x$ (because $x \oplus z \ominus x \overset{\text{commutativity}}{=} x \oplus -x \oplus z = z$). As G is finite, this y has to be unique, since otherwise there would not be enough elements left that can be added to x to result in all of the $|G| - 1$ possible remaining z'. Thus for any fixed x, z the probability that a randomly chosen y satisfies $x \oplus y = z$ is $\frac{1}{|G|}$.

Computing $\mathrm{lift}_{\mathrm{f}}^{\mathcal{G}}(S)$ consists of repeatedly adding group elements which by themselves are distributed uniformly at random if S was chosen randomly and f is a high quality hash function. Thus the accumulated value after every step is any given $z \in G$ with probability $\frac{1}{|G|}$, as is in particular the probability for the final result being equal to z. □

A more formal proof for xor specifically is given in section 6.2 of [MEL+21].

By the same argument, knowing the fingerprint for some set does not provide any information about the fingerprints for sets that differ by even only a single value. In conclusion, high quality fingerprints can be achieved by choosing any transitive and commutative operation for the monoid, for example xor or addition modulo 2^k, as long as values are mapped into the monoid with a high quality hash function.

While multiplication does not form a group when performed on the numbers in $[0, 2^k)_{\mathbb{N}}$, there still are groups based on multiplication modulo some number, e.g. $\mathbb{Z}_n^*$, the group yielded by multiplication modulo n on the set $\{x \in [0, n)_{\mathbb{N}} \mid x \text{ is coprime to } n\}$. In the following, when talking about multiplication, we will assume that the universe is chosen such that multiplication forms a group.

2.4 Cryptographically Secure Fingerprints

In the protocols for synchronizing data structures, fingerprints of sets are used for probabilistic equality checking: sets with equal fingerprints are assumed to be equal. Synchronization can thus become faulty if it involves unequal sets with equal fingerprints. If the universe of possible fingerprints is chosen large enough, and the distribution of fingerprints of randomly chosen sets is random within that universe, the probability for this to occur becomes negligible.

Random distribution of input sets is however a very strong assumption. What if a malicious party can influence the sets to be fingerprinted, with the goal of causing fingerprint collisions and consequently triggering faulty behavior of the system? Cryptographically secure fingerprints are an answer to this problem, being chosen such that it is computationally infeasible for an adversary to find inputs that lead to faulty synchronization.

2.4.1 General Considerations

A typical definition of cryptographically secure hash functions is the following [MVOV18]:

Definition 13. A *secure hash function* is a hash function $\mathrm{h} : U \to D$ that satisfies three additional properties:

pre-image resistance: Given $d \in D$, it is computationally infeasible to find a $u \in U$ such that $\mathrm{h}(u) = d$.

second pre-image resistance: Given $u \in U$, it is computationally infeasible to find a $u' \in U, u' \neq u$ such that $\mathrm{h}(u) = \mathrm{h}(u')$.

collision resistance: It is computationally infeasible to find $u, v \in U, u \ = v$ such that $\mathrm{h}(u) = \mathrm{h}(v)$.

Pre-image resistance has no influence on the vulnerability of the protocol to malicious actors, so all of the following discussion will focus on collision resistance only.

Since $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}(\{u\}) = \mathrm{f}(u)$, f must necessarily be collision resistant if $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$ is to be collision resistant. This alone is unfortunately not sufficient, we will see a specific counterexamples in section 2.4.2. Choosing a secure hash function f always comes with a performance cost, insecure hash functions usually take less time and less space to compute. If the synchronization protocol is only being run in a trusted environment, an insecure hash function might be preferable.

Whether a hash function is secure is not a binary dichotomy, but depends on what is considered "feasible" for an adversary. Greater security can usually be obtained at the cost of longer digests and longer computation times. Before presenting options for secure hash functions, we thus examine the impact of hash collisions first.

We can generally distinguish between malicious actors in two different positions: those who can actively impact the contents of the data structure to be synchronized, and those who passively relay updates and need to search for a collision within the available data. As a set of size n has 2^n subsets, if fingerprints are bit strings of length k, then by the pigeonhole principle a fingerprint collision can be found within any set of size at least $k + 1$.

An attack against the fingerprinting scheme by an active adversary can involve computing many fingerprints and adding the required items to the set once a collision has been found. Such an attack is not usable by the passive adversary. We will primarily focus on discussing active adversaries, as they are strictly more powerful than passive ones. Yet it should be kept in mind that passive adversaries can be more common in certain settings, particularly in peer-to-peer systems: if a node is interested in synchronizing a data structure, it probably trusts the source of the data, otherwise it would have little reason for expending resources on synchronization. The data may however be synchronized not with the original source but with completely untrusted nodes. Additionally, an active adversary that does not want to risk detection by adding suspicious items to the data structure is restricted to the operations of a passive adversary.

Fingerprint collisions result in parts of the data structure not being synchronized, so information is being withheld from one or both of the synchronizing nodes. When a malicious node synchronizes with an honest one, the malicious node can withhold

arbitrary information by simply pretending not to have certain data, which does not require finding collisions at all.

So the cases in which a malicious node can do actual damage by finding a collision are those where it supplies data to two different, honest nodes such that these two nodes perform faulty synchronization amongst each other. Specifically: let $\mathcal{M}$ be a malicious node, $\mathcal{A}$ and $\mathcal{B}$ be honest nodes, then a successful attack consists of $\mathcal{M}$ crafting sets X_A, X_B and sending these to $\mathcal{A}$ and $\mathcal{B}$ respectively, so that when $\mathcal{A}$ and $\mathcal{B}$ then run the synchronization protocol, they end up with different data structures. A passive adversary does not craft X_A, X_B but must find them as subsets of some set X supplied by an honest node. As we assume the underlying hash function f to be secure, at least one non-singleton set has to be involved in a collision.

There are some qualitative arguments that even if an adversary finds a fingerprint collision, the impact is rather low. Let $S_A \subseteq X_A$ and $S_B \subseteq X_B$ be nonequal sets with the same fingerprint. To have any impact on the correctness of a particular protocol run, their two fingerprints need to actually be compared during that run. For that to happen, they need to be of the form $[x_A, y_A)_{X_A}$ and $[x_B, y_B)_{X_B}$ respectively. The fingerprints of these ranges can then be compared if one of the nodes sends the fingerprint for the range $[\min(x_A, x_B), \max(y_A, y_B))_{X_i}$. That alone is still not sufficient, as any item within that range that is not part of the sets would change the fingerprint. So the two ranges actually need to be of the form $[\min(x_A, x_B), \max(y_A, y_B))_{X_A}$ and $[\min(x_A, x_B), \max(y_A, y_B))_{X_B}$, or simplified: there have to be $x, y \in U$ such that $S_A = [x, y)_{X_A}$ and $S_B = [x, y)_{X_B}$.

If the adversary has found such sets, that is still no guarantee that the range $[x, y)_{X_i}$ is being compared during the synchronization session of $\mathcal{A}$ and $\mathcal{B}$. For a set containing n items, there are $n^2 - (n-1) \in \mathcal{O}(n^2)$ distinct ranges, but only $\mathcal{O}(n)$ are compared in a given protocol run, since the worst-case message complexity is $\mathcal{O}(n)$ (see section 3.3.3). These numbers should only be considered as rough guidelines, they gloss over details such as the fact that there are more ranges containing roughly $\frac{n}{2}$ items then there are ranges containing almost all or almost no items. Yet they demonstrate that finding suitable colliding ranges still does not guarantee that a particular pair of nodes will be affected. In particular, there is no need for $\mathcal{A}$ and $\mathcal{B}$ to choose the range boundaries that occur in a protocol run deterministically (see section 3.3.4). They can even perform multiple randomized protocol runs in parallel, while keeping track of item transmissions and sending every item at most once across all these protocol runs.

Another factor mitigating the impact of an adversary finding fingerprint collisions is the communication with other, non-colluding nodes. A fourth party could send some $u \in U, x \preceq u \prec y$ to $\mathcal{A}$ or $\mathcal{B}$ before $\mathcal{A}$ and $\mathcal{B}$ synchronize, disrupting the collision.

Finally, in systems where nodes repeatedly synchronize with different other nodes, a single fingerprint collision in a single synchronization session would merely delay propagation of information rather than stop it completely (note that this does not hold for collisions of two singleton sets). Since the attack model requires both $\mathcal{A}$ and $\mathcal{B}$ to synchronize with more than one node in total, this might apply to many affected settings. Peer-to-peer systems communicating on a random overlay network in particular fall into this category. A malicious actor with enough control over the communication of other nodes to guarantee a tangible benefit from fingerprint collisions can likely disrupt operation of the network more effectively by exercising that

control than by sabotaging synchronization.

All of these arguments are however purely qualitative and should as such be taken into account with caution, they are not a substitute for quantitative cryptographic analysis. A strong attacker might be able to find many pairs of sets of colliding fingerprints, or many sets that all share the same fingerprint, and none of the above arguments consider these cases.

In a system with a consensus mechanism among all participating nodes, the choice of f can periodically be changed, with the frequency being some function of the time it takes to find a viable collision, the cost of rebuilding all auxiliary data structures, and the general level of paranoia among the participating nodes. The $\mathcal{O}(n)$ cost of rebuilding the data structures is then being amortized over the number of synchronization sessions occurring between rebuilds, which is still an improvement over protocols that need to perform $\mathcal{O}(n)$ computations per synchronization session.

2.4.2 Specific Monoids

Multiplication, addition and xor as ways of combining hashes have been studied in [BM97], in the context of hashing sequences, of which our ordered sets are a special case. Their setting requires not only associativity but also commutativity. They thoroughly break xor by reducing the problem of finding a collision to that of solving a system of linear equations. We will not restate the full attack, but we will describe a weaker but simpler mechanism based on similar ideas that allows finding subsets whose fingerprint is a specific target value. This can be used as an optimization in a trusted setting (see section 3.4.2).

The main observation is that xor is the additive operation in $\mathbb{F}_2$, the finite field on two elements $\{0, 1\}$. A fingerprint can be interpreted as a vector $\begin{bmatrix} b_1 & b_2 & \dots & b_k \end{bmatrix} \in \{0, 1\}^{1\times k}$. The fingerprint of a set $S = \{s_1, s_2, \dots, s_n\}$ can thus be computed as the sum (within $\mathbb{F}_2$, i.e. as the xor) of the vectors corresponding to $s_1, s_2, \dots, s_n$. The fingerprint of some $S' \subseteq S$ can also be regarded as the sum over all vectors, but each vector being first multiplied with a coefficient of 1 if $s_i \in S'$ and coefficient 0 if $s_i \notin S'$. In other words, the fingerprint of S' is a linear combination of the hashes of the items in S.

This enables us to efficiently find subsets whose fingerprint are a particular vector. Let $S = \{s_1, s_2, \dots, s_n\}$ be a set of items, and let $b = \begin{bmatrix} b_1 & b_2 & \dots & b_k \end{bmatrix} \in \{0, 1\}^{1\times k}$ be the target fingerprint. For $0 < i \leq n$, define $a_i = \begin{bmatrix} a_{i,1} & a_{i,2} & \dots & a_{i,k} \end{bmatrix} \in \{0, 1\}^{1\times k}$ to be $\mathrm{f}(s_i)$ interpreted as a vector over $\mathbb{F}_2$. The coefficients for linear combinations of the a_i equal to b are the solutions to the following system of linear equations over $\mathbb{F}_2$:

$$\begin{bmatrix} a_{1,1} & a_{1,2} & \dots & a_{1,k} \\ a_{2,1} & a_{2,2} & \dots & a_{2,k} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n,1} & a_{n,2} & \dots & a_{n,k} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

Solutions can be found by gaussian elimination in $\mathcal{O}(n^3)$ time, an exponential improvement over brute forcing by computing the fingerprints of all subsets.

As xor admits polynomial-time attacks by solving linear equations, the authors of [BM97] next consider addition and multiplication. They unify parts of

their discussion by relating the hardness of finding collisions to solving the balance problem: in a commutative group $(G, \oplus, \mathbb{0})$, given a set of group elements $S = \{s_1, s_2, \ldots, s_n\}$, find disjoint, nonempty subsets $S_0 = \{s_{0,0}, s_{0,1}, \ldots, s_{0,k}\} \subseteq S, S_1 = \{s_{1,0}, s_{1,1}, \ldots, s_{1,l}\} \subseteq S$ such that $s_{0,0} \oplus s_{0,1} \oplus \ldots \oplus s_{0,k} = s_{1,0} \oplus s_{1,1} \oplus \ldots \oplus s_{1,l}$. They then reduce the hardness of the balance problem to other problems.

For addition, the balance problem is as hard as subset sum, which was at the time of publication conjectured to be sufficiently hard. Wagner showed however in [Wag02] how to solve the balance problem in subexponential time for addition. To give an impression for the impact of this attack, consider [MGS15] which uses addition for combining SHA-3 [Dwo15] digests, producing fingerprints of length between 2688 and 4160 or 6528 to 16512 bits to achieve security levels of 128 or 256 bit respectively against Wagner's attack. [Lyu05] gives an improvement over Wagner's attack finding collisions in $\mathcal{O}(2^{n^\varepsilon})$ for arbitrary $\varepsilon < 1$, further weakening addition as a choice of monoid operation.

For multiplication, the balance problem is as hard as the discrete logarithm problem in the group. This is a more "traditional" hardness assumption than subset sum, there are groups for which no efficient algorithm is known. The main drawback is that multiplication is less efficient to compute that addition. [SMBA10] includes a comparison between the performance of addition and multiplication for incremental hashing, the additive hash outperforms the multiplicative one by two orders of magnitude, even though it uses longer digests to account for Wagner's attack.

For our context in which fingerprints are frequently sent over the network, longer computation time might still be preferable over longer hashes. Fingerprints based on multiplication nevertheless need larger digests than traditional, non-incremental hash functions. [MSTA17] suggests fingerprints of 3200 bit to achieve 128 bit security, and motivated by this uses binary elliptic curves as the underlying group to achieve more compact fingerprints - fingerprints of length $2k$ bits give $\mathcal{O}(2^k)$ security.

[BM97] also proposes a fourth monoid based on lattices. [LKMW19] give a specific instantiation providing 200 bits of security with fingerprints of size $16 \cdot 1024 = 16384$ bit.

All of the preceeding monoid operations are commutative, which our approach to fingerprint computation does not require. A typical associative but not commutative operation is matrix multiplication. Study of a family of hash functions based on multiplication of invertible matrices was initiated in [Zém91], whose security is related to solving hard graph problems on the Cayley graph of the matrix multiplication group. [PQ+11] gives a good overview about the general principle and the security aspects of Cayley hash functions.

While [TZ94], an improvement over the originally proposed scheme, has been successfully attacked in [GIMS11] [PQ10], there are several modifications [Pet09] [BSV17] [Sos16] for which no attacks are currently known, and [MT16] showed random self-reducibility for Cayley hash functions.

Whereas the schemes based on commutative groups operate on two bitstrings at a time, the Cayley hash functions operate on two individual bits at a time. Attacks thus assume that the manipulated bits can be freely chosen, using this to e.g. craft palindromic inputs. This may mean that such attacks are hard to apply to our setting where the input bits are produced in non-reorderable batches by another, cryptographically secure hash function. After finding a bit sequence that yields a Cayley hash collision, an attacker still needs to find items for which the

concatenation of their hashes is the desired it sequence.

Alternatively, one can forgo the additional hash function and apply the Cayley hash function directly to the encodings of items. This simplifies the overall scheme and removes reliance on a second hash function. On the other hand, if one expects the additional hash function to be harder to attack than the Cayley hash, then making the attacks against the Cayley hash function less flexible might overall make attacks more difficult. Furthermore, Cayley hash functions are usually slower than typical choices for the hash function from items to bitstrings, so using both can produce significant speedups if items have long encodings.

Aside from Cayley hashes, we are not aware of any non-commutative monoids used for hashing. We would further like to note that any hashes based on groups still have more structure than required for our designs, as we don't need existence of inverse elements. Strictly speaking our designs do not even require a neutral element: the fingerprint of the empty set never needs to be exchanged in a protocol run, so we could have just as well defined fingerprints for all nonempty sets without relying on the neutral element to do so. The presentation we chose is merely more elegant. And finally, we do not require pre-image resistance for our fingerprints. Suitable hash functions could thus be located in a more general design space than studied in any literature we know of.

2.5 Pseudorandom Data Structures

Monoidal fingerprints allow efficient deterministic computation of fingerprints, but at the price of having to rely on some non-standard cryptographic primitives. A more established cryptographic construction is that of Merkle trees [Mer89] and related data structures, all based on hashing the concatenation of child hashes. There is a significant body of work in the context of authenticated data structures based on these constructions, from rather simple balanced binary trees [NN00] to arbitrary acyclic directed graphs [MND+04], and we will be able to refer to such work for proofs of collision resistance.

The Merkle construction is however not associative when using a cryptographically secure hash function h, as $\mathrm{h}(\mathrm{h}(\mathrm{h}(a) \cdot \mathrm{h}(b)) \cdot \mathrm{h}(c)) = \mathrm{h}(\mathrm{h}(a) \cdot \mathrm{h}(\mathrm{h}(b) \cdot \mathrm{h}(c)))$ would constitute a violation of collision resistance. Different such tree representations of the same set thus lead to completely different fingerprints. So in order to use this construction, there must be a unique search tree representation defined for every set.

[Sny77] shows that deterministic search trees with unique shapes require $\mathcal{O}(\sqrt{n})$ time for insertions and deletions in the worst case, making them unsuitable for our use case. A natural extension is to look at randomized data structures that define a unique presentation which allows modification and search in $\mathcal{O}(\log(n))$ with high probability.

This has been thoroughly researched in the field of *history-independent data structures*, intuitively speaking data structures whose bitlevel representations do not leak information about their construction sequence. This requirement has been shown in [HHM+05] to be equivalent to having a unique representation. For more background, we refer to the excellent introduction of [BBJ+16].

In the following, we present two such data structures for representing sets - a randomized binary tree, and the skip list - and define Merkle-like fingerprinting

schemes for them such that the fingerprints of arbitrary ranges can be computed in logarithmic time with high probability. As all nodes need to arrive at the same representation, no true randomness can be involved. Instead, all random decisions are based on pseudorandom bits derived from the data itself.

These schemes allow cryptographically secure fingerprints whose collision resistance does not depend on uncommon cryptographic building blocks. The price to pay is that computing fingerprints can take linear time in the worst case. Note however that the communication complexity of any synchronization protocol using these fingerprints is completely unaffected by the randomized computation time.

Unfortunately, adversaries can efficiently create data sets for which the randomized solutions degraded to linear performance, thus enabling denial-of-service attacks as the cost of computing the fingerprints during a protocol run can be made to dominate the communication cost. Passive adversaries however cannot effectively attack these fingerprints: even though they can technically only forward those parts of the data structure that have degraded performance relative to the number of items in those parts, they would still increase the absolute resource usage of their victims if they simply forwarded all data.

Proponents of randomized data structures also claim that they are significantly easier to implement than balanced search trees (see e.g. [SA96] or [Pug90]), making it worthwhile to consider them even in trusted environments, in which their expected logarithmic update complexities suffice.

2.5.1 Pseudorandom Treaps

A *treap* [SA96] is a search tree in which every vertex has an associated *priority* from a totally ordered set, and in which the priorities of the vertices form a heap:

Definition 14. Let U be a set, $\prec$ a linear order on U, P a set, $\leq$ a linear order on P, priority : $U \to P$, $V \subseteq U$, and T a binary search tree on V.

Then we call T a treap if for all $v \in V$ with parent p it holds that $\text{priority}(v) \leq \text{priority}(p)$.

From the properties of treaps shown in [SA96] we will rely on the following:

- If the priorities of all vertices are pairwise unequal, then there is exactly one treap on the vertex set, which is equal to of the search tree obtained by successively inserting items in decreasing priority without any balancing.
- If the priorities are distributed uniformly at random, the height of the treap is expected to be logarithmic in the number of vertices.
- Inserting or deleting an item while maintaining the treap properties can be done in time proportional to the height of the treap.

In the following, we will fix some cryptographically secure hash function $\mathrm{p} : U \to \{0,1\}^k$, and define $\text{priority}(u) := \mathrm{p}(u)$, using the numeric order on $\{0,1\}^k$ for comparing priorities. Since p is collision resistant, we can assume the resulting treaps to be unique, and since the output of a secure hash function is indistinguishable from a random mapping, we can assume the resulting treaps to have expected logarithmic height.

Treaps store items in every vertex, whereas Merkle trees only store items in their leaves. [BLL02] proposes the following natural generalization and proves that the labels are collision free if the underlying hash function h is collision free:

$$\mathrm{tlbl_h}(v) := \begin{cases} \mathrm{h}(v), & \text{for leaf } v \\ \mathrm{tlbl_h}(\mathrm{tlbl_h}(c_<) \cdot \mathrm{tlbl_h}(v)) & \text{v internal vertex with left child } c_< \\ & \text{and no right child} \\ \mathrm{tlbl_h}(\mathrm{tlbl_h}(v) \cdot \mathrm{tlbl_h}(c_>)) & \text{v internal vertex with right child } c_> \\ & \text{and no left child} \\ \mathrm{tlbl_h}(\mathrm{tlbl_h}(c_<) \cdot \mathrm{h}(v) \cdot \mathrm{tlbl_h}(c_>)) & \text{v internal vertex with left child } c_< \\ & \text{and right child } c_> \end{cases}$$

We will call a treap labeled according to this scheme a *Merkle treap*. Since Merkle treaps are unique, we can define the fingerprint of some set as the root label of the Merkle treap on that set if it is nonempty, or `Nothing` otherwise.

Subsets

Now that we have defined the fingerprint of a set, we need a way of efficiently computing the fingerprint of any range $[x, y)_S$ given the Merkle treap of the set S. The algorithm can be expressed in exactly the same framework as that of section 2.2.1, changing merely some type definitions and the `combine` function:

```
-- A fingerprint is now either a hash or Nothing.
type Fingerprint = Maybe D

emptyFp :: Fingerprint
emptyFp = Nothing

liftFp :: D -> Fingerprint
liftFp d = Just d

-- Combine three fingerprints into a new fingerprint.
combine :: Fingerprint -> Fingerprint -> Fingerprint -> Fingerprint
combine Nothing Nothing Nothing = Nothing
combine (Just a) Nothing Nothing = Just a
combine Nothing (Just b) Nothing = Just b
combine Nothing Nothing (Just c) = Just c
combine (Just a) (Just b) Nothing = Just (hash (concat a b))
combine (Just a) Nothing (Just c) = Just (hash (concat a c))
combine Nothing (Just b) (Just c) = Just (hash (concat b c))
combine (Just a) (Just b) (Just c) = Just (hash (concat (concat a b) c))
```

While the tree traversal and accumulation of data otherwise stays the same, the nature of the correctness argument changes quite a bit, since we cannot rely on the associativity of `combine` anymore. Instead, the core idea of the argument is that the computed fingerprint depends only on the relative positioning between the tree vertices within the range, which stays identical no matter how many vertices outside the range are added.

Recall that the shape of a treap is always that of the tree obtained by inserting items sorted by decreasing priority without rebalancing. Now consider the vertex v with the smallest distance to the root r such that $x \preceq v \prec y$, i.e. the vertex computed by `findInitial`. All items within $[x, y)_S$ are contained in $\mathrm{V}(T_v)$. The fingerprint computation for the range is clearly local to T_v, and while insertion of items outside that tree might change the path from r to v, they do not change the shape of T_v.

It thus remains to show that `accGeq` and `accLt` compute the same result no matter how many items outside the range are contained in the tree they process. We give the full argument for `accGeq`, the case for `accLt` follows analogously.

Proof. Let v be the root of a tree T_v and $x \in U$. We show by induction on $|\mathrm{V}(T_v)|$ that `accGeq` computes a value equal to the root label of the Merkle treap on $S_{\succeq} := \{u \in \mathrm{V}(T_v) \mid x \preceq u\}$ (or `Nothing` if the set is empty).

IB: If v is a leaf with $v \prec x$, then $(\texttt{accGeq v x}) = \texttt{Nothing}$ and indeed $S_t = \emptyset$. If v is a leaf with $x \preceq v$, then $(\texttt{accGeq v x}) = \texttt{Just}\,\mathrm{h}(\texttt{v}) = \texttt{Just}\,\mathrm{tlbl_h}(\texttt{v})$.

IH: Let $n \in \mathbb{N}$ and T_r be a tree rooted at r with $|\mathrm{V}(T_r)| = n$. Then $(\texttt{accGeq v x}) = \mathrm{tlbl_h}(v)$.

IS: Let T_v be a tree rooted at v with $|\mathrm{V}(T_v)| = n + 1$, let $c_<$ be the right child of v and let $c_>$ be the left child of v.

Case 1. If $v \prec x$, then:

$$\begin{aligned}(\texttt{accGeq v x}) &\overset{def}{=} (\texttt{accGeq c}_> \texttt{ x})\\ &\overset{IH}{=} \mathrm{tlbl_h}(c_>)\\ &\overset{v \prec x}{=} \mathit{tlabel}\mathrm{h}(v).\end{aligned}$$

Case 2. Otherwise, $x \preceq v$. Neither v nor any vertex in $T_{c_<}$ has any influence on the shape of $T_{c_>}$. We thus have:

$$\begin{aligned}(\texttt{accGeq v x}) &\overset{def}{=} (\texttt{combine (accGeq c}_< \texttt{ x) h(v)}\,\mathrm{tlbl_h}(\texttt{c}_>))\\ &\overset{IH}{=} (\texttt{combine}\ \mathrm{tlbl_h}(\texttt{c}_<)\ \mathrm{h}(\texttt{v})\,\mathrm{tlbl_h}(\texttt{c}_>))\\ &\overset{def}{=} \mathrm{tlbl_h}(\mathrm{tlbl_h}(c_<) \cdot \mathrm{h}(v) \cdot \mathrm{tlbl_h}(c_>))\\ &\overset{def}{=} \mathrm{tlbl_h}(v).\end{aligned}$$

The computation time is proportional to the height of the treap, so expected $\mathcal{O}(\log(n))$. But as the hash function is not associative, the tail-recursive versions of `accGeq` and `accLt` do not compute the correct result. Intuitively, the problem is that the computation traverses the tree from top to bottom, but the fingerprints are defined through a computation from the bottom to the top. Accordingly, tail-recursive formulations can be achieved by adding parent-pointers to all vertices and performing the computation by first traversing to the range boundary and then accumulating fingerprints while moving back up through the tree. The space complexity thus remains in $\mathcal{O}(1)$.

Just as with balanced search trees, a binary treap might be the most natural formulation, but not the most efficient one on real hardware. [Gol09] describes B-treaps, which are treaps of higher degree. If such an optimized realization is chosen, the need to map between the actual data layout and the shape of the treap which defines the fingerprints complicates the implementation, in particular compared to monoidal fingerprints which naturally abstract over tree implementation details.

Beyond treaps there are other randomized search trees. [PT89] suggests hash tries for fingerprinting sets, and our mechanism for computing fingerprints of ranges is compatible with their approach. The reason we opted for treaps is that the pseudorandom selection of the tree shape is completely decoupled from the ordering over the items, whereas in a hash trie the ordering follows from the tree shape (or vice versa, depending on the viewpoint). As our protocols can benefit from a non-arbitrary ordering (compare section 3.1.1), treaps are the more flexible solution.

There exists also some work on deterministic unique tree representations that stay efficient if the number of items is either a particularly small or a particularly large fraction of the size of the universe, see [ST94] for both a specific approach and more related work. These approaches could also be adapted for our purposes, but since they pose restrictions on the number of items that can be stored, they are less flexible than our other approaches.

Adversarial Treap Construction

The main motivation for looking at pseudorandom schemes was to be able to rely on well-known, conventional, non-associative hash functions, implying that resilience against malicious peers is desired. While Merkle-treaps are collision-resistant, they open up another angle of attack: a malicious party can try to construct unbalanced treaps to make the cost of computing fingerprints super-logarithmic in the number of items, leading to a denial-of-service attack as the cost of computing the fingerprints during a protocol around dominates the communication cost.

There is a fairly straightforward way of creating treaps on n vertices of height n in expected $\mathcal{O}(n^2)$ time and $\mathcal{O}(1)$ space, making treaps unsuitable for an adversarial environment. A sequence $(u_0, u_1, \ldots, u_{n-1})$ of items such that the treap on these items has height n can be computed by successively choosing arbitrary but always increasing with respect to $\preceq$ items p_j and letting $u_i := p_j$ if $\lfloor i \cdot \frac{2^k}{n} \rfloor \leq \text{priority}(p_j) < \lfloor (i+1) \cdot \frac{2^k}{n} \rfloor$ (recall that k is the number of bits of a hash digest). The probability for the priority of an arbitrary item to fall within the desired range is $\frac{1}{n}$, so finding one takes expected $\mathcal{O}(n)$ time, finding a sequence of n of these accordingly takes expected $\mathcal{O}(n^2)$ time.

2.5.2 Pseudorandom Skip Lists

The skip list [Pug90] [Pug98] is a probabilistic data structure for storing sets that does not use a tree representation, but rather a hierarchy of progressively sparser linked lists containing set items in sorted order. The list at layer 0 contains every item, the list at layer $i+1$ contains every item from layer i with a probability of $\frac{1}{2}$. Intuitively, when inserting an item into the skip list, one performs coin flips until one loses, the item is then added at the correct position into as many layers as one performed coin flips. Figure 2.2 gives an example.

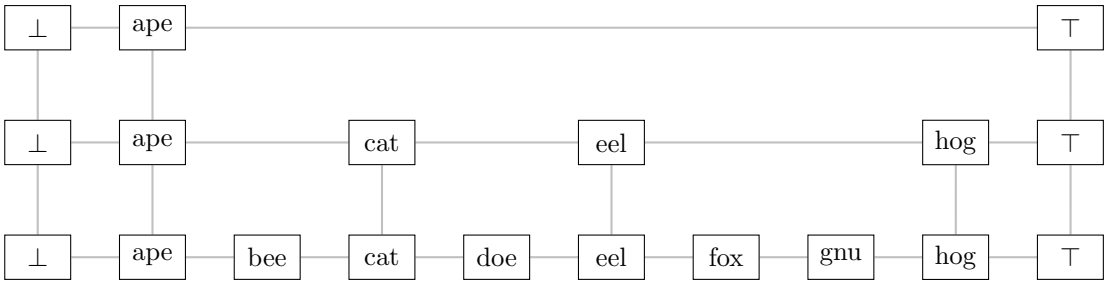

Figure 2.2: An example skip list, demonstrating the layers of sorted linked lists. The $\perp$ and $\top$ vertices are the start and end points for all lists.

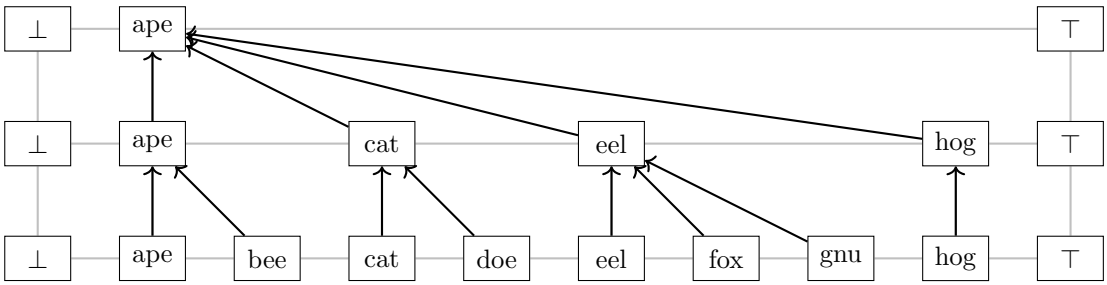

Figure 2.3: A visualization of which labels influence which other labels. One can see that these dependencies form a tree.

To define a pseudorandom skip list, we again fix some cryptographically secure hash function $\mathrm{p} : U \to \{0,1\}^k$, and determine the layers of an item u by counting the number of leading zero bits of $\mathrm{p}(u)$ and adding one.

Toward a fingerprinting scheme, we define a labeling over all items in all layers. Let $u_{l,i}$ denote the i-th item in layer l, recall that h is some cryptographically secure hash function. Then $\mathrm{slbl_h}(u_{0,i}) := \mathrm{h}(u_{0,i})$ and $\mathrm{slbl_h}(u_{l+1,i}) := \mathrm{h}(\mathrm{slbl_h}(u_{l,j}) \cdot \mathrm{h}(\mathrm{slbl_h}(u_{l,j+1}) \cdot \mathrm{h}(\ldots \cdot \mathrm{slbl_h}(u_{l,j+m}))))$ where $(u_{l,j} \ldots u_{l,j+m})$ is the longest subsequence of layer l such that $u_{l,j} = u_{l+1,i}$ and $u_{l,j+m} \prec u_{l+1,i+1}$. Figure 2.3 visualizes an example.

Every label contributes to the computation of at most one other label, which is located on the next layer. Maintaining labels when inserting or deleting an item thus takes a number of hash computations bounded by the height of the skip list, which is expected $\mathcal{O}(\log(n))$. The number of label concatenations that need to be performed at each layer is also randomized, but on average there are two, as the probability for an item to also occur at the next layer is $\frac{1}{2}$. The expected time thus stays in $\mathcal{O}(\log(n))$.

Denote by $\mathrm{domain}(u_{l,i})$ the set of items u such that $u_{l,i} \preceq u \prec u_{l,i+1}$, i.e. the set of all items whose hash contributes to the label of $u_{l,i}$. Now let $S \subseteq U$ be a finite set. The fingerprint of S is defined as `Nothing` if $S = \emptyset$, and otherwise as $\mathrm{h}(u_{l_0,i_0} \cdot \mathrm{h}(u_{l_1,i_1} \cdot \mathrm{h}(\ldots \cdot u_{l_m,i_m})))$ where $(u_{l_0,i_0}, u_{l_1,i_1}, \ldots, u_{l_m,i_m})$ is the unique shortest sequence of vertices in the pseudorandom skip list on S such that its items are strictly increasing and their domains partition S, choosing the vertex of the highest possible layer if multiple vertices have equal domains. See fig. 2.4 for an example.

We next sketch an algorithmic way of determining this sequence for any range $[x, y)_S$ given the deterministic skip list for S, but refrain from giving correctness proofs as they are rather verbose and convey no particular insight that is not better

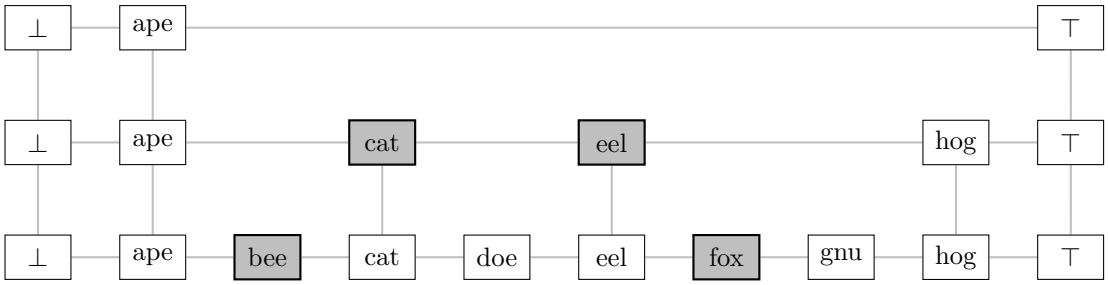

Figure 2.4: The skip list nodes contributing to the fingerprint of the range $[\text{bee}, \text{gnu})_S$. If the range had been $[\text{bee}, \text{fox})_S$, it would have again been the topmost eel vertex whose label would have been hashed.

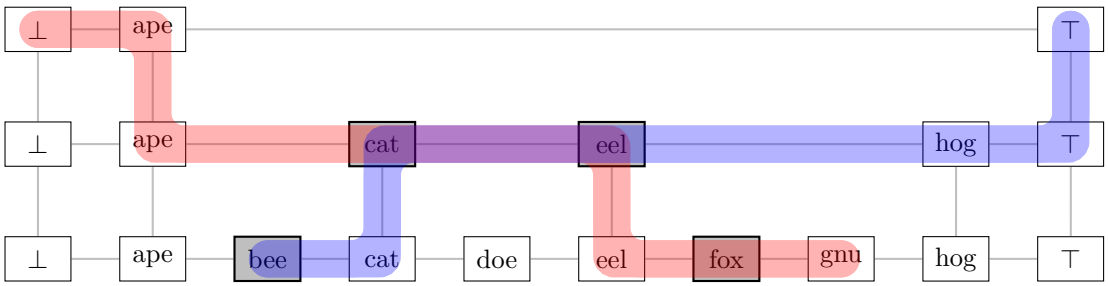

Figure 2.5: The forward search path to gnu (starting at the top-left corner) and the backward search path to bee (starting at the top-right corner) allow to read off which vertices contribute to the fingerprint of $[\text{bee}, \text{gnu})_S$.

conveyed in the accompanying figures.

Let $\bar{u} = (u_0, u_1, \ldots, u_m)$ be the shortest path in the skip list from its beginning to the vertex y excluding y itself, i.e. the sequence of vertices traversed when looking up y in the skip list. This sequence has expected logarithmic length and can be computed in expected logarithmic time. Obtain a shorter sequence $\bar{u}' = (u'_0, u'_1, \ldots, u'_{m'})$ by keeping from all vertices in $\bar{u}$ that correspond to the same item the one in the greatest layer such that its domain does not contain any item greater than or equal to y. Figure 2.5 provides an example.

Now let $\bar{v} = (v_0, v_1, \ldots, v'_o)$ be the shortest path in the skip list obtained by inverting the order on items from its beginning to the vertex x. Intuitively, this corresponds to looking up x "from the back" in a doubly-linked skip list. Obtain from $\bar{v}$ the corresponding shorter sequence $\bar{v}' = (v'_0, v'_1, \ldots, v'_{o'})$ by keeping from all vertices that correspond to the same item the one of the greatest layer such that its domain does not contain any item strictly less than x.

We claim that $\bar{v}'$ and $\bar{u}'$ always intersect (if $[x, y)_S \neq \emptyset$) in at least one vertex, i.e. there are unique $v'_i \in \bar{v}'$ and $u'_j \in \bar{u}'$ with $v'_i = u'_j$. Then the sequence $(v'_0, v'_1, \ldots, v'_i = u'_j, u'_{j+1}, \ldots, u'_m)$ is the unique shortest sequence of vertices of strictly increasing items whose domains partition $[x, y)_S$. Since $\bar{v}'$ and $\bar{u}'$ can be computed in expected logarithmic time, so can this sequence, and thus fingerprints can be computed in expected logarithmic time.

Adversarial Skip List Construction

Pseudorandom skip lists are even less suited for adversarial environments then pseudorandom treaps. Performance degrades if most items have the same maximum layer. A randomly chosen item has a hash beginning with a 1 bit with probability $\frac{1}{2}$, these items only reside in layer 0. An adversary can thus create a linear list of exclusively layer 0 items of length n in expected $2n \in \mathcal{O}(n)$ time.

Chapter 3

Set Reconciliation

In this chapter, we consider the set reconciliation protocol sketched in the introduction in greater detail. We define the protocol in section 3.1, prove its correctness in section 3.2, and do a complexity analysis in section 3.3. Section 3.4 lists some optimizations which do not change the asymptotic complexity but which avoid some unnecessary work. We conclude the chapter with an example application in section 3.5, briefly describing how the protocol can be applied to the synchronization of the hash graphs that arise, e.g., in the context of distributed version control systems such as git [CS14].

3.1 Recursive Set Reconciliation

The set reconciliation protocol assumes that there is a set U, a linear order $\preceq$ on U, a node $\mathcal{X}_0$ locally holding some $X_0 \subseteq U$, and a node $\mathcal{X}_1$ locally holding $X_1 \subseteq U$. Both nodes are able to compute fingerprints for sets via some function $\mathrm{fp} : \mathcal{P}(U) \to D$. This function could in principle be any hash function, but in the complexity analysis we will assume that it is instantiated as $\mathrm{lift}_{\mathrm{f}}^{\mathcal{M}}$ with a suitable underlying hash function f and a suitable monoid $\mathcal{M}$ over the universe D (see definition 6).

Recall that $[x, y)_S$ denotes the set $\{s \in S \mid x \preceq s \prec y\}$ if $x \prec y$, or $S \setminus [y, x)_S$ if $y \preceq x$ (definition 8).

$\mathcal{X}_0$ and $\mathcal{X}_1$ exchange messages over multiple rounds, a message consists of an arbitrary number of *range fingerprints* and *range item sets*. A range fingerprint is a triplet $(x, y, \mathrm{fp}([x, y)_{X_i}))$ for $x, y \in U$, a range item set a four-tuple (x, y, S, b) for $x, y \in U$, $S \subseteq [x, y)_{X_i}$, and $b \in \{0, 1\}$. b indicates whether the range item set is a response to a previous range item set.

When a node $\mathcal{X}_i$ receives a message, it performs the following actions:

- For every range item set (x, y, S, b) in the message, all items in S are added to the locally stored set X_i. If $b = 0$, the node then adds the range item set $(x, y, [x, y)_{X_i} \setminus S, 1)$ to the response, unless $[x, y)_{X_i} \setminus S = \emptyset$.

- For every range fingerprint $(x, y, \mathrm{fp}([x, y)_{X_j}))$ in the message, it does one of following:

  Case 1 (Equal Fingerprints). If $\mathrm{fp}([x, y)_{X_j}) = \mathrm{fp}([x, y)_{X_i})$, nothing happens.

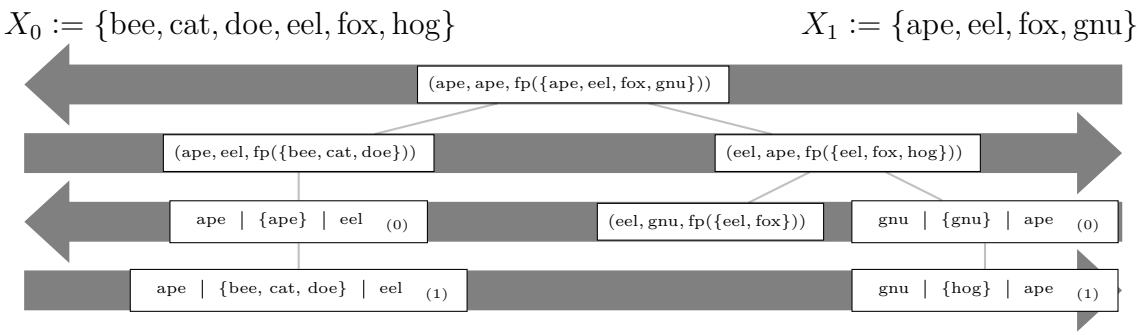

Figure 3.1: An example run of the protocol. $\mathcal{X}_1$ initiates reconciliation for all items between ape and ape(ordered alphabetically) by sending its fingerprint for the whole range.
Upon receiving this range fingerprint, $\mathcal{X}_0$ locally computes $\text{fp}([\text{ape}, \text{ape})_{X_0})$. Since the result does not match the received range, $\mathcal{X}_0$ splits X_0 into two parts of equal size and transmits range fingerprints for these subranges.
In the third round, $\mathcal{X}_1$ locally computes fingerprints for the two received ranges, but neither matches. $|[\text{ape}, \text{eel})_{X_1}| \leq 1$, so $\mathcal{X}_1$ transmits the corresponding range items set, i.e. $(\text{ape}, \text{eel}, \text{ape}, 0)$. $|[\text{eel}, \text{ape})_{X_1}| > 1$, so another recursion step is performed. After splitting the range, the lower range is large enough to send its fingerprint, the upper one however only contains one item and thus results in another range item set.
In the fourth and final communication round, $\mathcal{X}_0$ receives two range item sets and answers with the items it holds within those ranges. When it receives the range fingerprint $(\text{eel}, \text{gnu}, \text{fp}([\text{eel}, \text{gnu})_{X_1}))$, it computes an equal fingerprint for $(\text{eel}, \text{gnu}, \text{fp}([\text{eel}, \text{gnu})_{X_0}))$, so no further action is required for this particular range.

Case 2 (Recursion Anchor). The node may add the range item set $(x, y, [x, y)_{X_i} 0)$ to the response. If $|[x, y)_{X_j}| \leq 1$, it must do so.

Case 3 (Recurse). Otherwise, the node selects $m_0 = x \prec m_1 \prec \ldots \prec m_k = y \in U$, $k \geq 2$ such that among all $[m_0, m_1)_{X_i}$ for $0 \leq l < k$ at least two ranges are non-empty. For all $0 \leq l < k$ it adds either the range fingerprint $(m_l, m_{l+1}, \text{fp}([m_l, m_{l+1})_{X_i}))$ or the range item set $(m_l, m_{l+1}, [m_l, m_{l+1})_{X_i} 0)$ to the response.

- If the accumulated response is nonempty, it is sent to the other node. Otherwise, the protocol has terminated successfully.

To initiate reconciliation of a range $[x, y)$, a node $\mathcal{X}_i$ sends a message containing solely the range fingerprint $(x, y, \text{fp}([x, y)_{X_i}))$.

Figure 3.1 gives an example run of the protocol.

3.1.1 Observations

Partitioning based on a total order allows the nodes to perform a limited form of queries, i.e. range queries. A node can ask for reconciliation within a certain range, rather than over the whole universe.

If the universe U is not finite, then there are items that require an arbitrary amount of bytes to encode. Since the protocol needs to transmit items to denote

range boundaries, no reasonably complexity guarantees can be given for infinite universes. We will thus assume U to be finite and small enough that items can be reasonably encoded. This assumption is not very restrictive in practice because nodes can always synchronize hashes of items rather than the items themselves. The protocol can then be either followed by a phase where hashes of interest are transferred and anwered by the actual items, or the protocol can be made aware of the distinction and use hashes as range boundaries while transmitting actual items for range item sets.

When reconciling hashes in place of actual items, any semantically interesting order on the items would be replaced by an arbitrary order on the hashes. But rather than using only's the hashes as range boundaries, one can just as well add additional information. For example if the universe of interest consists of timestamped strings of arbitrary length, the range boundaries can consist of timestamped hashes, ordered by timestamp first and using the numeric order on the hashes as a tiebreaker. Section 3.5 gives a more detailed example for utilizing this technique.

3.2 Proof of Correctness

We now prove the correctness of the protocol. The protocol is correct if for all $x, y \in U$ both nodes eventually hold $[x, y)_{X_i} \cup [x, y)_{X_j}$ after a node $\mathcal{X}_i$ has received a message pertaining to the range $[x, y)$.

Case 1 (Range Item Set). If the message contains the range item set $(x, y, [x, y)_{X_j} 0)$, then $\mathcal{X}_i$ adds all items to its set, resulting in $[x, y)_{X_i} \cup [x, y)_{X_j}$ as desired. $\mathcal{X}_j$ then receives $(x, y, [x, y)_{X_i} \setminus [x, y)_{X_j}, 1)$, ending up with $[x, y)_{X_j} \cup ([x, y)_{X_i} \setminus [x, y)_{X_j}) = [x, y)_{X_i} \cup [x, y)_{X_j}$ as desired.

Case 2 (Range Fingerprint). Otherwise, the message contains a range fingerprint $(x, y, \mathrm{fp}([x, y)_{X_j}))$.

Case 2.1 (Equal Fingerprints). If $\mathrm{fp}([x, y)_{X_j}) = \mathrm{fp}([x, y)_{X_i})$, the protocol terminates immediately and no changes are performed by any node. Assuming no fingerprint collision occurred, $[x, y)_{X_i} = [x, y)_{X_j} = [x, y)_{X_i} \cup [x, y)_{X_j}$ as desired.

Case 2.2 (Recursion Anchor). If $\mathcal{X}_i$ adds the range item set $(x, y, [x, y)_{X_i} 0)$, then case 1 applies when the other node receives the response, with the roles of the nodes reversed.

Case 2.3 (Recurse). Let $count_i := |[x, y)_{X_i}|$ and $count_j := |[x, y)_{X_j}|$. $count_j \geq 2$, since otherwise $\mathcal{X}_j$ would have sent an item set for the range. Similarly, $count_i \geq 2$, since we are not in case 2.2. Thus, $count_i + count_j \geq 4$, and the protocol has already been proven correct for all cases where $count_i + count_j < 4$.
We can thus finish the proof by induction on $count_i + count_j$, using the induction hypothesis that for all $x', y' \in U$ such that $|[x', y')_{X_i}| + |[x', y')_{X_j}| < count_i + count_j$ the protocol correctly reconciles $[x', y')_{X_i}$ and $[x', y')_{X_j}$.
$\mathcal{X}_i$ partitions the range into $k \geq 2$ subranges, of which at least two must be nonempty. Thus $|[m_l, m_{l+1})_{X_i}| < count_i$ for all $0 \leq l < k$. Furthermore, $[m_l, m_{l+1})_{X_j} \subseteq [x, y)_{X_j}$ and thus $|[m_l, m_{l+1})_{X_j}| \leq |[x, y)_{X_j}|$, so overall we have $|[m_l, m_{l+1})_{X_i}| + |[m_l, m_{l+1})_{X_j}| < count_i + count_j$ and can apply the induction hypothesis to conclude that every subrange is correctly reconciled. Since the subranges partition the original range, the original range is then correctly reconciled as well.

3.3 Complexity Analysis

The protocol gives nodes the freedom to respond to a range fingerprint with a range item set even if the range fingerprint is arbitrarily large. For a meaningful complexity analysis we need to restrict the behavior of the node, a realistic modus operandi is for a node to send a range item set whenever it holds a number of items less than or equal to some threshold $t \in \mathbb{N}, t \geq 1$ within the range. Higher choices for t reduce the number of roundtrips, but increase the probability that a items is being sent even though the other node already holds it.

A node is similarly given freedom over the number of subranges into which to split a range when recursing. We will assume a node always splits into at most $b \in \mathbb{N}, b \geq 2$ subranges. As with t, higher numbers reduce the number of roundtrips at the cost of potentially sending items or fingerprints that did not need sending.

Because we want to analyze not only the worst-case complexity but also the complexity depending on the similarity between the two sets held by the participating nodes, we define some more fine-grained instance size parameters: n_0 and n_1 denote the number of items held by $\mathcal{X}_0$ and $\mathcal{X}_1$ respectively. We let $n := n_0 + n_1$, $n_{\min} := \min(n_0, n_1)$ and $n_\triangle := |([x, y)_{X_0} \cup [x, y)_{X_1}) \setminus ([x, y)_{X_0} \cap [x, y)_{X_1})|$.

3.3.1 Preliminary Observations

A helpful observation for the following analysis is that the range fingerprints that are being exchanged during a protocol run form a rooted tree where every vertex has at most b children. When a leaf of the tree is reached, an exchange of range item sets follows. Equal fingerprints can also cut the tree short, but for the following worst-case analyses we will assume this does not occur.

Node $\mathcal{X}_i$ can branch at most $\lceil \log_b(n_i) \rceil$ times, so the overall height of the tree is bounded by $2 \cdot \lceil \log_b(n_{\min}) \rceil$. The number of vertices of such a complete tree of height h is at most $\sum_{i=0}^{h} b^i = \frac{b^h - 1}{b-1}$. For $h \leq 2 \cdot \lceil \log_b(n_{\min}) \rceil$, $\frac{b^h-1}{b-1} \leq 2 \cdot 2 \cdot n_{\min} \leq 2n \in \mathcal{O}(n)$.

The parameter t determines when recursion is cut off, and thus influences the height of the tree. For $t = 1$, the protocol recurses as far as possible. For $t = b$, the last level of recursion is cut off, for $t = b^2$ the last two levels, and so on. Overall, the height of the tree is reduced by $\lfloor \log_b(t) \rfloor$.

3.3.2 Communication Rounds

The number of communication rounds clearly corresponds to the height of the tree, plus 2 to account for the exchange of range item sets, so the worst-case is $2 + 2 \cdot \lceil \log_b(n_{\min}) \rceil - \lfloor \log_b(t) \rfloor \in \mathcal{O}(\log_b(n))$. This number cannot be bounded by $n_\triangle$, as witnessed by problem instances where one node is missing exactly one item compared to the other node. In such an instance, $b - 1$ branches in each recursion step result in equal fingerprints, but the one branch that does continue reaches the recursion anchor only after the full number of rounds. See fig. 3.2 for an example.

3.3.3 Communication Complexity

The total number of bits that needs to be transmitted during a protocol run is proportional to the number of vertices in the tree. Every range fingerprint consists of two items and one fingerprint, so assuming U is finite this lies in $\mathcal{O}(1)$. Since

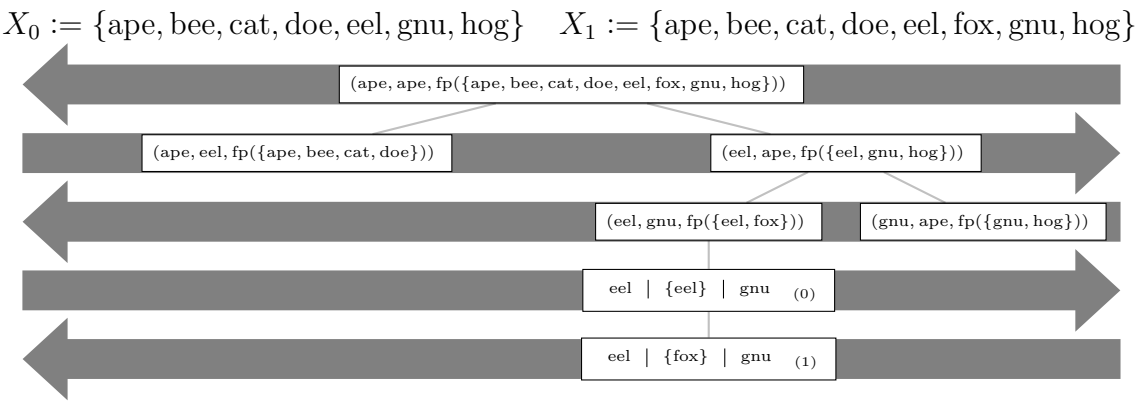

Figure 3.2: An example run of the protocol that takes the greatest possible number of rounds even though $n_\triangle = 1$. $b := 2, t := 1$.

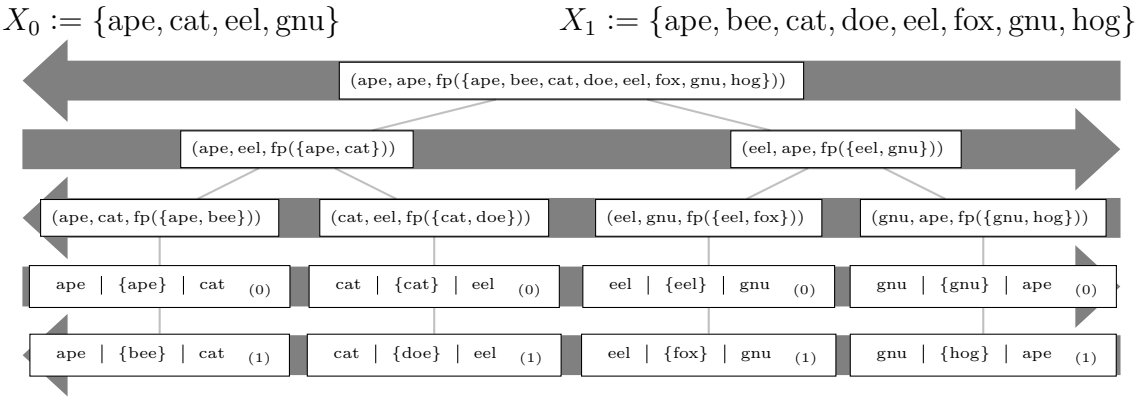

Figure 3.3: An example run of the protocol that requires transmitting the maximum amount of bytes. $b := 2, t := 1$.

there are at most $2n$ vertices in the tree, the range fingerprints require at most $\mathcal{O}(n)$ bits to be communicated.

The exchange of range item sets consists in the worst case of exchanging every item using $\lceil \frac{n}{t} \rceil$ range item sets. A range item set needs to transmit two items to encode the boundaries, as well as the items themselves, which lies in $\mathcal{O}(1)$ per range item set. All range item sets together thus amount to another $\mathcal{O}(n)$, leading to a total of $\mathcal{O}(n)$ bits being transmitted in the worst case.

Figure 3.3 shows a worst-case example in which the tree of height $h := \log_b(2 \cdot n_{\min})$ has all $\frac{b^h - 1}{b-1}$ vertices.

This worst-case complexity is no better than the naïve protocol where the two nodes simply exchange their whole item sets. The motivation behind our protocol is to utilize similarities in the sets that are being reconciled, so it it makes sense to bound the communication complexity by $n_\triangle$.

As discussed in the analysis of the number of communication rounds, if $n_\triangle = 1$ this results in the protocol tracing a path from the root to a leaf in the tree. If $n_\triangle$ takes on larger values, more paths from the root to leaves occur, each contributing up to $2 + 2 \cdot \lceil \log_b(n_{\min}) \rceil - \lfloor \log_b(t) \rfloor \in \mathcal{O}(\log_b(n))$ vertices, but the number of vertices in the union of these paths is of course bounded by the total number of vertices in the tree. We thus obtain a communication complexity of $\mathcal{O}(\min(n_\triangle \cdot \log_b(n), n))$ bits.

3.3.4 Impact of Randomized Range Boundaries

Section 2.4.1 argues in favor of randomizing range boundaries to protect against the impact of maliciously crafted fingerprint collisions. We now consider three different randomization strategies and their impacts on the number of communication rounds and the number of bits that need to be transmitted. To describe the boundary selection strategies, first consider that splitting an ordered set $S = \{s_0 \prec s_1 \prec \ldots \prec s_{n-1}\}$ items into b ranges consists of choosing b indices $i_0 < i_1 < \ldots < i_{b-1}$, thus defining the ranges $[s_{i_0}, s_{i_1})_S, [s_{i_1}, s_{i_2})_S, \ldots, [s_{i_{b-2}}, s_{i_{b-1}})_S, [s_{i_{b-1}}, s_{i_0})_S$.

The randomization strategy with the lowest impact creates ranges of identical sizes as the regular strategy (splitting into b subranges of evenly distributed size per communication around), but shifts where these boundaries lie. Precisely: first, range boundaries $i_0, i_1, \ldots, i_{b-1}$ are computed deterministically, all of which lie $\frac{n}{b}$ apart. Then some random integer $k \in [0, \frac{n}{b})_{\mathbb{N}}$ is chosen, and then the new boundaries are $s_{i_j+k \mod b}$ for all $0 \leq j < b$.

As the number and size of ranges is identical to that of the deterministic protocol version, communication complexity remains unchanged. Yet there are $\frac{n}{b}$ different possible partions of a range of size n, thus reducing the probability that a particular maliciously crafted fingerprint collision affects a given communication round by that factor.

In order to minimize the probability that a collision affects a protocol run, range boundaries can be chosen fully at random, for a total of $\binom{n}{b}$ possible outcomes. The worst-case number of communication rounds occurs when in every round $b-1$ of the range contain exactly one item and the last range contains all the remaining items, leading to $\frac{n}{b-1}$ rounds. The expected number of rounds is however still $\mathcal{O}(\log_b(n))$ with high probability, as it corresponds to the height of a randomly chosen b-complete tree, which is expected to be in $\mathcal{O}(\log_b(n))$ [Dev90].

A simple way of recovering a logarithmic worst-case round number while keeping the number of possible partions super-linear is to choose the boundaries randomly within equally sized regions of the range: choose i_j randomly from $[j \cdot \frac{n}{b}, (j+1) \cdot \frac{n}{b})_{\mathbb{N}}$ (and take i_{b-1} modulo n if necessary). The number of possible outcomes is $(\frac{n}{b})^b \in \mathcal{O}(n^b)$. Since any range obtained this way can contain most $2b$ items, the worst-case number of communication rounds is $\mathcal{O}(\log_{\frac{b}{2}}(n))$.

Finally note that if reconciliation is initiated by one node sending the fingerprint over all its items, no randomization is involved in this first round, which could be utilized by the malicious actor crafting fingerprint collisions. To protect against this, the first reconciliation round should already start with sending b (randomized) range fingerprints. Or in other words, a node wishing to initiate reconciliation simply acts as if it had just received a range fingerprint over the whole universe not matching the fingerprint over all its items. In addition to protecting against fingerprint collisions, this also reduces the number of communication rounds by one, though at the cost of having sent an unnecessary amount of data if both nodes are already holding the same set of items.

3.3.5 Computational Complexity

We now analyze the computational cost incurred by a single communication round, i.e. computing the response to a message. This includes both fingerprint comparisons as well as locating the items to transmit. We do however assume that an

auxiliary data structure is available to help with this computation, e.g. the fingerprint tree structure presented in section 2.2. We exclude both space usage and maintenance cost for this data structure from the per-round complexity.

We will assume that transferring an item as part of a range item set requires $\mathcal{O}(1)$ time and space. The relevant computational overhead per communication round thus consists of computing $\mathrm{fp}([x, y)_{X_i})$ for every received range fingerprint $(x, y, \mathrm{fp}([x, y)_{X_j}))$, as well as partitioning $[x, y)_{X_i}$ in case of a mismatch and computing the fingerprints over all subranges. These computations can be performed independently for all received range fingerprints, so in particular they can be performed sequentially, reusing space. The overall space complexity of the per-round computations is equal to that of the computations for a single range fingerprint.

We first address the problem of partitioning ranges. After the indices of the boundaries have been determined as described in the previous section, how can they be efficiently translated into actual items? This can be accomplished by augmenting the auxiliary tree structure by storing in each vertex v the number of vertices in T_v. With this information, one can efficiently traverse the tree based on relative offsets on the number of items to traverse.

The problem thus consists of computing $k \cdot (b+1)$ fingerprints in the worst case, where k is the number of range fingerprints that have been received - one fingerprint computation per received fingerprint to compare the received fingerprint with the local one, and then b further ones for splitting the range.

A naive approach is to query the auxiliary data structure for each fingerprint to compute individually. The maximum amount of queries is upper-bounded by n, as n items cannot be split into more than n non-empty ranges. The greatest possible number of range fingerprints within a single round corresponds to the number of leaves in the recursion tree, which is in $\mathcal{O}(n)$. The auxiliary data structure requires $\mathcal{O}(\log(n))$ time per query, leading to an overall time complexity of $\mathcal{O}(n \cdot \log(n))$.

Instead, the bulk fingerprint computation from section 2.2.2 should be used, leading to a time complexity of $\mathcal{O}(\min(k \cdot \log(|S|), |S|))$.

3.3.6 Discussion

Some observations beyond the pure statement of the complexities are worth pointing out.

First, the choice of branching factor b and the item size k at which range item sets are being sent give a lot of flexibility to the protocol. For $b = 2$ and $k = 1$ the protocol minimizes the work spent on items that lie outside the symmetric difference between the two stored sets, at the cost of more "meta traffic" that needs to be transmitted. In contrast, the trivial reconciliation protocol in which both nodes immediately send all their items minimizes this meta traffic by instead transmitting items whose transmission would not have been necessary. Increasing the choices of b and k effectively interpolates between these two extremes. They can thus be chosen based on which resource is more valuable, bandwidth or time.

The second remark is that range-based reconciliation makes good use of the "free complexity budget" inherent to the problem: storing a set of items in sorted order for efficient access is usually done via balanced trees, implying a storage overhead of $\mathcal{O}(n)$ and insertion/deletion time complexity of $\mathcal{O}(\log(n))$. The auxiliary tree structure for computing hashes requires exactly the same resources (and can in fact be

implemented by storing additional data in the vertices of the storage tree) and thus does not increase the overall asymptotic complexity at all. Solutions without such auxiliary data structures are in some sense wasting parts of the available complexity budget.

Intuitively, an ideal solution should have communication complexity proportional to $n_\triangle$. The range-based approach does not achieve this goal, but it misses it by merely a logarithmic factor. The weakest aspect is the logarithmic number of communication rounds, considering that the trivial protocol easily achieves a constant number. On the plus side, increasing b and/or k can counteract this up to the point where the increased bandwidth outweighs the reduced number of communication rounds.

3.4 Smaller Optimizations

We now give a list of optimizations which do not impact the overall complexity analysis, but which do improve on some constant factors.

3.4.1 Non-Uniform Partitions

When partitioning a range into subranges, the protocol does not specify where exactly to place the boundaries. Splitting into partitions of roughly equal sizes makes a lot of sense if new data could arise anywhere within the linear order with equal probability. Is however the items are likely to fall within certain ranges of the order, it can be more efficient to use more fine-grained partitions within those regions. If for example items are sorted by timestamp, and new items are expected to be propagated to every node in the system within a couple of seconds, then all items older than ten seconds can be comfortably lumped together in a large range.

3.4.2 Subset Checks

When a node $\mathcal{X}_i$ receives a range fingerprint $(x, y, \mathrm{fp}([x, y)_{X_j}))$, it might have a different fingerprint for that range, but one of the resulting subranges could match the received fingerprint. The node can then ignore that range, transmitting only the remaining ones.

This scenario is a special case of receiving the fingerprint of a subset of the items one holds within a range. In principle, a node can compute the fingerprints of arbitrary subsets of its range, trying to find a match. If a match is found, the node knows that it holds a superset of the items the other node holds within the range. The protocol could be extended with a message part that transmits items and does not warrant a response, this would be used to transmit $[x, y)_{X_i} \setminus [x, y)_{X_j}$.

Computing the fingerprint of all subsets of a range is infeasible since there are 2^n many subsets. As discussed in section 2.4.2 however, if the group operation for fingerprint computation is xor, a node can check whether the fingerprint of a subset of items matches a given fingerprint in $\mathcal{O}(n^3)$. Assuming no fingerprint collisions occurred, checking whether a subset of $[x, y)_{X_i}$ matches $\mathrm{fp}([x, y)_{X_j})$ is equivalent to checking whether $[x, y)_{X_i}$ is a superset of $[x, y)_{X_j}$.

This leads to a protocol with a per-round computational complexity of $\mathcal{O}(n^4)$[1] ($\mathcal{O}(n)$ fingerprint computations of complexity $\mathcal{O}(n^3)$ each) and the same worst-case guarantees as the basic protocol, but which can skip recursive steps more often than the basic protocol. In particular, whenever two nodes reconcile where one of them holds a subset of the items of the other node, reconciliation terminates after at most three communication rounds.

3.4.3 Efficiently Encoding Ranges

A naive encoding of range fingerprints and range item sets would transmit both boundaries of every range, transmitting $2n$ items if there are n ranges in a message. This can be brought down to $1+n$ by utilizing that the lower boundary of all but the first range is also the upper boundary of the preceding one. An efficiently encoded message consists of the number of ranges it contains, followed by the lower boundary of the first range, followed by pairs of range information and the upper boundary of the range that the information pertains to, where each *range information* is either a fingerprint, a set of items, or a dummy value signaling that this part of the overall range is already fully synchronized.

The knowledge that range boundaries are transmitted in ascending order can be used for compression. As a simple example, if range boundaries are natural numbers, every upper boundary can be transmitted as the difference to the previous boundary. This way, smaller numbers are transmitted, which admit shorter encodings. The same argument applies for transmitting sets of items in sorted order.

Note additionally that two adjacent range item sets can be merged into a single one, saving on the transmission of the boundary between them.

3.4.4 Utilizing Range Boundaries

Whenever the lower boundary of a range is transmitted, an actual item is transmitted. If the receiving node knew whether the other node held this item, it would automatically be reconciled and could be excluded from further recursion steps. One way for achieving this is to tag each range with a bit to indicate whether the lower boundary is being held by the sender. A different approach is to require that ranges are only split at items which the splitting node holds, this way no explicit bits needs to be transmitted, yet many boundaries can be identified as being held by the other node.

Recall that we assume U to be finite. For $x \in U$ we can thus denote by successor(x) the necessarily unique $z \in U$ such that $x \prec z$ and there exists no $y \in U$ such that $x \prec y \prec z$ if a z with $u \prec z$ exists, or $\min(U)$ otherwise. When sending or receiving a range from x to y for which both nodes know that the sending node holds x, both nodes act as if the range went from successor(x) to y, except that the receiving node adds x to its local set of items.

[1] This can be brought down to $\mathcal{O}(n^3)$ by using techniques for solving multiple systems of linear equations based on the same matrix more efficiently, e.g. by precomputing a LU decomposition or a QR decomposition.

3.4.5 Multi-Level Fingerprints

Since the protocol interprets equal range fingerprints as both nodes holding the same set of items within the range, but never verifies this merely probabilistic assumption, fingerprints need to be long enough to guarantee low collision probability. The longer the fingerprints, the more bits need to be transmitted however.

One can use smaller fingerprints if in case of equal fingerprints the nodes then exchange an additional fingerprint rather than immediately stopping the recursion for this range. Whenever a node receives a first-level range fingerprint that is equal to its own, it answers with the second-level fingerprint for the same range. When a node receives a second-level range fingerprint that is equal to its own, it terminates the recursion. If it is not equal, it recurses as usual, in particular it sends first-level fingerprints for any large subranges.

This scheme can of course be extended to an arbitrary number of fingerprint levels. Every level increases the worst-case number of roundtrips by one, but decreases the average message size. If cryptographically secure fingerprints are desired, low levels of the fingerprint hierarchy do not have to be secure, only the top level fingerprint needs to be so. Alternatively, each fingerprint level can consist of some substring of a long, cryptographically secure fingerprint, such that the concatenation of these substrings yields the original fingerprint.

3.5 Reconciling Hash Graphs

In this section, we demonstrate how to apply the set reconciliation protocol the problem of reconciling hash graphs. Hash graphs arise in contexts where pieces of data refer to other pieces of data by a secure hash computed over the data to be referred to. Distributed version control systems such as git [CS14] for example represent the evolving contents of a directory as a set of deltas which describe how the contents changed from an earlier version, this earlier version is referenced as the hash of another such object. Some objects describe how to merge conflicting concurrent changes, these objects can reference multiple other objects. Since addressing uses a secure hash function, an object can only reference objects which existed prior to it. All in all, objects can thus form arbitrary directed, acyclic graphs (dags).

When two nodes wish to synchronize, they update each other's histories to the union of all objects known to both of them. This is a natural setting for the set reconciliation protocol. Since objects can become arbitrarily large, one would reconcile sets of hashes of these objects, and then in a second stage request the actual objects for all newly obtained hashes.

This approach completely ignores the edges of the dag. While at first glance it might seem inefficient to ignore available structural information, the fact that the dags can take arbitrary shapes makes it hard to utilize the edge structure. The graphs might be dense or sparse, could contain arbitrarily large independent sets, paths, tournaments etc.

Even though arbitrary dags are possible and thus the logarithmically bounded worst-case complexity is important, one can in practice make some reasonable assumptions about the hash graphs arising from an average, version-controlled repository. It seems likely for example that most concurrent new work is performed relative to a rather recent state of the repository, whereas work based off a very old

state is rather unlikely. We could hope for better average reconciliation times if the reconciliation algorithm leveraged this expectation.

To that end, rather than reconciling merely hashes of objects, we can reconcile pairs $(\mathit{depth}, \mathit{hash})$, sorting by depth first and hash second. The *depth* of an object is the length of the longest path from that object to a root object, i.e., an object without predecessors. That is, the depth of an object is 1 greater than the greatest depth of any predecessor object. If most concurrent work is based off similarly new state, then it also falls into a similar range of the linear order. A protocol run does not need to recurse into ranges of low depth then, since no concurrently working peer produces new objects of low depth. If this assumption turns out to be wrong, the protocol nevertheless upholds its worst-case guarantee of a logarithmic number of communication rounds.

Another common scenario is that a node has not produced any new local objects and merely wants to catch up with the current state. Let d be the maximal depth among all objects the node holds. Then rather than reconciling the range $[(0,0),(0,0))$, the node can reconcile the range $[(0,0),(d+1,0))$ as well as sending an empty range item set for the range $[(d+1,0),(0,0))$. All new changes of depth $d+1$ or greater are then fetched in only two communication rounds. Any concurrent work (from a depth perspective) is reconciled as usual, and if all new changes are based off a recent version, the whole reconciliation process only takes a single round trip.

Chapter 4

Protocol Variants

In this chapter, we sketch some adaptations of the set reconciliation algorithm to related problems. Section 4.1 considers more general partitioning strategies than one-dimensional ranges. Section 4.2 presents how to reconcile key-value mappings. Section 4.3 shows an asymmetric variant of the protocol, mirroring the contents of a set or map held at a primary node to a replica node. Section 4.4 demonstrates how our data structures for fingerprint computations can also serve as simple authenticated data structures.

4.1 Higher-Dimensional Ranges

The set reconciliation algorithm recursively partitions the item sets of both nodes by splitting the sets into successive ranges. But for the correctness of the algorithm, the exact mechanism by which partitions are chosen is irrelevant. All that is necessary is that in each communication round, when receiving a mismatching fingerprint for some non-empty $S \subseteq U$, the node chooses a partition $S_0 \dot\cup S_1 \dot\cup \ldots \dot\cup S_k = S$ of S, and then communicates to the other node the choice of the partition and how its items are distributed among the individual sets.

Partitioning into successive ranges over a linearly ordered set has some nice properties: It can be applied to virtually any set since one can always arbitrarily define a linear order. The choice of the partition and the way the items are distributed among the sets can be encoded and transmitted efficiently. And finally, when receiving an arbitrarily chosen partition into successive ranges, a node can always efficiently compute the items it holds within that partition as well as the fingerprint over these items.

As long as nodes always wants to synchronize the complete sets they hold, this is completely sufficient. But more elaborate partitioning strategies can become attractive if nodes might wish to only synchronize certain subsets. The partitioning mechanism effectively dictates which such subsets can be described, taking on the role of a query language. A natural next step is to look for more powerful partitioning mechanisms that still admit efficient fingerprint computation.

We now give one such partitioning scheme by generalizing the one-dimensional ranges of the original presentation of the algorithm to k-dimensional ranges. Let U be a finite set of items, and let $\preceq_0, \preceq_1, \ldots, \preceq_{k-1}$ be linear orders over U. A *k-dimensional range* is a k-tuple $((x_0, y_0), (x_1, y_1), \ldots, (x_{k-1}, y_{k-1}))$ with $x_i \preceq_i y_i$. Given some $S \subseteq U$, the items from S in this range are all items $v \in S$ such that

$x_i \preceq_i v \preceq_i y_i$ for all $0 \leq i < k$. Note that for $k = 1$ this is equivalent to the ranges used in previous chapters. The reconciliation protocol needs to be adapted so that messages include the k-dimensional range boundaries, but otherwise no changes are required.

Fingerprints can be computed efficiently by storing the set as a (balanced) k-d tree [Ben75]. As k-d tree are binary trees, the labels from chapter 2 can be used without any modification. Fingerprint computation still works by traversing the tree, alternating between the dimensions just as in e.g. a k-d tree item lookup. For $k = 1$, this yields exactly the algorithms of chapter 2.

Generalized partitioning schemes, including k-dimensional ranges, can not only be applied to the regular set reconciliation protocol, but also to all modifications discussed in the further sections.

4.2 Map Reconciliation

A key-value mapping, i.e. a partial function $\mathrm{m} : K \rightharpoonup V$ with a finite domain from some set of keys K to a set of values V, can be reconciled by reconciling the set $\{(k, v) \mid k \in \mathrm{dom}(\mathrm{m}) \text{ and } v = \mathrm{m}(k)\}$. After this reconciliation, a node may have obtained two pairs $(k, v), (k, v')$ if the two nodes mapped the same key k to distinct values v, v', so the resulting set would not correspond to an updated map. In those cases, both nodes compute the single new image of k as $\mathrm{f}(v, v')$, where $\mathrm{f} : V \times V \to V$ is some function known to all participating nodes.

Particularly interesting are cases where f can be computed via another interactive protocol. If for example V consists of the fingerprints of finite subsets of some universe U, and $\mathrm{f}(v, v')$ is defined to be the fingerprint of the union of the two sets whose fingerprints are v and v', then the two nodes can run a set reconciliation session to efficiently obtain the union. Viewn more abstractly, when reconciling a map, the values can be reconciled via arbitrary nested protocol invocations.

4.3 Set and Map Mirroring

We based the presentation of our synchronization approach on set reconciliation because it is a symmetric problem where both node use identical algorithms to compute identical types of messages. A related, asymmetric problem is that of a *replica* node setting its locally stored set to that of a *primary* node, utilizing similarity between the two initial sets to minimize the communication complexity. The approach of recursively exchanging fingerprints for subsets of decreasing size can be modified to solve this *set mirroring* problem.

The primary node can run exactly the same protocol as that for such reconciliation. The replica node uses a slightly modified version. Whenever it receives a range item set, it adds the received items to its local set as usual, but then it deletes all item it holds within that range which were not part of the received range item set. Whenever the replica node sends a range item set, it sends an empty one.

The complexity analysis is identical to that of the set reconciliation protocol, the correctness argument is analogous: ranges with equal fingerprints are already correctly mirrored if no fingerprint collisions occurred, exchanging range item sets

correctly mirrors all items within that range, and large ranges with non-equal fingerprints can be handled recursively because mirroring the partitions of a set results in mirroring the whole set.

Just as for reconciliation, maps can be mirrored by interpreting them as sets of key-value pairs. Mirroring maps has some interesting use cases, for example filesystems can be regarded as maps from paths to strings. The efficient creation of backups then becomes the problem of mirroring such a map onto a backup server that may already hold a similar, older backup. An equivalent problem is that of efficiently distributing source code updates from a server to clients which may hold old versions of the source code.

4.4 Authenticated Data Structures

Authenticated data structures solve the problem of outsourcing data structure membership queries processing to untrusted replicas rather than processing all queries at the original trusted data source. The trusted source publishes a short digest to a client. The client can then send a query to a replica, which answers with the result and a small certificate which together with the digest proves that the answer is indeed correct. For more details on this three-party model and pointers to the rich literature on the topic, we refer to [MND+04].

Merkle trees form a simple authenticated set. The root label is the digest, and the certificate for an affirmative membership query consists of the labels of the children of the vertices on the path from the root to the item in question. These labels can be used to recompute the root label, proving that the item is part of the original tree. Fabricating a sequence of labels to fake inclusion of an item amounts to breaking the hash function. Non-membership queries can be authenticated by providing certificates for tree membership of items stored in adjacent leaves such that one is less than and one is greater than the item in question.

Our examination of randomized data structures was driven by the need for Merkle-like properties, so it is not very surprising that they can also be used as authenticated data structures. Indeed [NN00] mentions treaps as an alternative to regular Merkle trees. Monoidal labels based on Cayley hash functions remove the need for a specific tree shape to be maintained. Skip lists as authenticated data structures have been studied in [GT00], but they require the hash function to be commutative. Our construction removes this assumption.

Chapter 5

Related Work

The range-based approach to set reconciliation has been employed as an implementation detail for a specific use case in [CEG+99] section 3.6 and [SYW+17] section II.A[1]. Neither work studies it as a self-contained approach of interest to the set reconciliation problem, consequently the treatment is rather superficial, e.g,. neither work considers collision resistance, even though it is relevant to both settings. Considering the simplicity of the approach, there may well be further publications reusing or reinventing it. To the best of our knowledge, there is no prior work dedicated to the algorithm itself.

We give an overview of the existing literature on set fingerprinting in section 5.1, and of the set reconciliation literature in section 5.2.

5.1 Fingerprinting Sets

As range-based synchronization relies on fingerprinting sets, we give a small overview of work on that topic. Some of these references have already come up withn this thesis because they are directly related to our approaches, the other references are given to round out the picture and because they might prove useful for developing further ideas.

Merkle trees [Mer89] introduce the idea of maintaining hashes in a tree to allow efficient updating without a group structure on the hashes. Since the exact shape of the tree determines the root label, unique representations of sets as trees are of interest. [Sny77] proves the important negative result that in general unique representations require superlogarithmic time to maintain under insertions and deletions. [ST94] gives logarithmic solutions for sparse or dense sets and points to further deterministic solutions.

Unique representations with probabilistic time complexities that are logarithmic in the expected case have been studied in [PT89], suggesting hash-tries as a set representation for computing hashes. Pugh also developed skip lists [Pug90], a probabilistic set representation not based on trees. [SA96] offers treaps, another randomized tree structure. Further study of uniquely represented data structures has been done under the moniker of *history-independent data structures*, including

[1]We cite a survey because there is no standalone publication on CCNx 0.8 Sync. The survey refers to online documentation at `https://github.com/ProjectCCNx/ccnx/blob/master/doc/technical/SynchronizationProtocol.txt`

IO-efficient solutions for treaps [Gol09] and skip list [BBJ+16] in an *external memory* computational model.

Beyond the comparison of root labels for set equality testing, Merkle trees and their technique of hashing the concatenation of hashes form the basis of many authenticated data structures, ranging from simple balanced trees or treaps [NN00] and skip lists [GT00] to more general DAG-based data structures [MND+04].

A different line of authenticated data structures utilizes dynamic accumulators [CL02], small digests for sets that can be efficiently updated under insertion and deletion of items, and which allow computing small certificates that some item has been ingested. [PTT16] and [PTT11] use accumulators to provide authenticated set data structures that are more efficient than their Merkle-tree-based counterparts. Accumulators are stronger than necessary for the range-based synchronization approach, so they are not discussed in our main text. An interesting opportunity for further research could be the design of synchronization protocols which leverage the strong properties offered by dynamic accumulators.

Orthogonal to these lines of research are algebraic approaches. The earliest mention of hashing into a group and performing computations on hashes to compactly store information about sets that we could find is in [WC81]. They provide probabilistic set equality testing, but without maintaining any tree structure, effectively anticipating (cryptographically insecure) homomorphic set fingerprinting.

The first investigation of the cryptographic security of this approach was done in the seminal [BM97], albeit in the slightly less natural context of sequences rather than sets, with the additive hash being broken in [Wag02] and [Lyu05]. The multiset homomorphic formulation was given in [CDVD+03], further constructions are given in [CNQ09], [MSTA17] and [LKMW19].

The generalized hash tree of [PSTY13] extends the idea of combining hashes by a binary operation to operations that are not closed, instead the output is mapped back into the original domain by a separate function. These "compressed" outputs are then arranged in a tree, this technique allows using the algebraic approach for authenticated data structures.

5.2 Set Reconcilliation

We now present several alternative approaches to the set reconciliation problem, almost all of which perform reconciliation in a single communication round at the price of high computational cost. In the following presentations, we consider the setting of a node $\mathcal{X}_0$ holding a set $X_0 \subseteq U$ reconciling with a node $\mathcal{X}_1$ holding a set $X_1 \subseteq U$. We denote by $n_\triangle$ the size of the symmetric difference between X_0 and X_1.

5.2.1 Characteristic Polynomial Interpolation

The seminal work on set reconciliation is [MTZ03], introducing an approach based on characteristic polynomial interpolation (CPI). They consider the items to reconcile as bit-strings of length b. Given a set $S = \{x_1, x_2, \ldots, x_n\}$, the characteristic polynomial $\chi_S(Z)$ is defined as $(Z - x_1)(Z - x_2) \ldots (Z - x_n)$. The elements in S are exactly the roots of $\chi_S(Z)$, so given the coefficients of $\chi_S(Z)$, S can be recovered by factoring the polynomial.

Next observe that when dividing one characteristic polynomial by another, factors occurring in both polynomials cancel each other out, so it holds that:

$$\frac{\chi_{X_0}(Z)}{\chi_{X_1}(Z)} = \frac{\chi_{X_0 \setminus X_1}(Z)}{\chi_{X_1 \setminus X_0}(Z)}$$

The degrees of $\chi_{X_0 \setminus X_1}(Z)$ and $\chi_{X_1 \setminus X_0}(Z)$ are upper-bounded by $n_\triangle$. Assume that the nodes have somehow obtained an approximation $\bar{n}_\triangle \geq n_\triangle$. The nodes then agree on $\bar{n}_\triangle$ sample points from some sufficiently large field, and node $\mathcal{X}_i$ locally evaluates $\chi_{X_i}(Z)$. The results are exchanged, so each node can then recover the value of $\frac{\chi_{X_0 \setminus X_1}(Z)}{\chi_{X_1 \setminus X_0}(Z)}$ at each of the sample points.

Knowing these values, a node can interpolate them to recover $\chi_{X_0 \setminus X_1}(Z)$ and $\chi_{X_1 \setminus X_0}(Z)$. Factoring these then yields the items $X_0 \setminus X_1$ and $X_1 \setminus X_0$, completing the reconciliation.

The number of evaluation results that need to be exchanged is linear in $n_\triangle$, thus the total number of bits that need to be transmitted is proportional to $n_\triangle$, which is more efficient than the range-based approach. Evaluating $\chi_{X_i}(Z)$ at $\bar{n}_\triangle$ sample points takes $\mathcal{O}(|X_i| \cdot \bar{n}_\triangle)$ time, since every item in X_i increases the number of terms in the characteristic polynomial. So while the communication complexity is asymptotically optimal, the computational complexity for a node is at least linear in the size of its set.

Interpolating the evaluation points can be reduced to performing Gaussian elimination in $\mathcal{O}(\bar{n}_\triangle^3)$ time. The authors also give an $\mathcal{O}(\bar{n}_\triangle^3)$ time algorithm for finding the roots of the interpolated polynomials. The overall computational complexity is thus in $\mathcal{O}(|X_i| \cdot \bar{n}_\triangle + \bar{n}_\triangle^3)$. The space complexity of the computations is not elaborated on, but since recovering of the polynomials requires the information about all $\bar{n}_\triangle$ evaluation points, it appears to be at least in $\mathcal{O}(\bar{n}_\triangle)$, i.e. an implementation using only a fixed amount of memory is not possible.

The above description of the protocol assumes that an upper bound $\bar{n}_\triangle \geq n_\triangle$ is known, it then performs reconciliation within a single round trip. For the protocol to be asymptotically efficient in the communication complexity, $\bar{n}_\triangle$ must be within a constant factor of $n_\triangle$. If no such value can be known in advance, the number of sample points can be increased by a constant factor over multiple communication rounds, resulting in an appropriate $\bar{n}_\triangle$ after a logarithmic number of rounds, leading to an overall $\mathcal{O}(\log(\bar{n}_\triangle))$ communication rounds. In each round, the nodes verify probabilistically whether $\bar{n}_\triangle$ is large enough by sending some additional random sample points and verifying whether the result of the interpolation correctly evaluates these samples. A linear number of verification sample points decreases the probability of false positives exponentially.

5.2.2 Invertible Bloom Lookup Tables

A bloom filter [Blo70] is a probabilistic data structure for set membership queries. Given an item, the bloom filter can answer a membership query with a certain false positive rate. The size of a bloom filter is a fraction of the size of an array storing the items. The larger the filter is chosen relative to the number of items it stores, the smaller is the false positive rate.

Invertible bloom lookup tables [GM11] (IBLTs) extend this behavior with the ability to list all items stored in the data structure, albeit again with an error

probability which decreases as the fraction of space per contained item goes up. [EGUV11] introduces set reconciliation based on IBLTs by formulating a way to compute the difference between two IBLTs, with the result corresponding to an IBLT encoding the corresponding set difference. This leads to the following set reconciliation algorithm:

Assume the nodes have obtained an approximation $\bar{n}_\triangle \geq |X_0 \setminus X_1|$. $\mathcal{X}_0$ creates an IBLT large enough that listing its contents succeeds with sufficient probability as long as it contains at most $\bar{n}_\triangle$ items. It then stores all items of X_0 in this table and transmits the table to $\mathcal{X}_1$. $\mathcal{X}_1$ creates an IBLT of the same size containing the items of X_1, and then subtracts the tables, yielding a table containing most $\bar{n}_\triangle$ items. This table is invertible, so $\mathcal{X}_1$ can recover $X_0 \setminus X_1$ from it. The same protocol can run in parallel with the roles reversed, resulting in full reconciliation.

Creating the IBLT requires $\mathcal{O}(|X_i|)$ time for node $\mathcal{X}_i$ and $\mathcal{O}(\bar{n}_\triangle)$ space, both of which dominate the cost for subtracting the IBLTs and for inverting them.

Similar to the CPI approach, the nodes need an estimate of the size of the set difference before they can perform reconciliation. The authors present a single-message estimation protocol for the set difference that uses IBLTs as well. The size of the message is in $\mathcal{O}(\log(|U|))$. Both creating or receiving and using the message requires $\mathcal{O}(X_i)$ time for node $\mathcal{X}_i$ and $\mathcal{O}(\log(|U|))$ space.

Overall, the IBLT approach achieves set reconciliation in a single round trip, transmitting only $\mathcal{O}(n_\triangle + \log(|U|))$ bits. The computational cost is however linear in the size of the sets, and the space requirements for the computation are in $\mathcal{O}(n_\triangle)$, which can cause trouble the case of lopsided reconciliation sessions.

5.2.3 Other Bloom Filter Approaches

Several other approaches based on bloom filters have been published, all of which achieve a constant number of roundtrips and a small message size at the cost of at least linear computation time and computation space requirements. [BCM02] transmits the nodes of a patricia tree representation of the sets in a bloom filter, [TZX+11] uses a bloom filter-based approach to obtain an estimate of the size of the set difference in the first phase of a CPI reconciliation.

[GL12] employs counting bloom filters, [LGR+19] makes use of cuckoo filters. [OAL+19] combines IPLTs with regular bloom filters to reduce the message size.

5.2.4 Partition Reconciliation

[MT02] is the only publication we found which considers reconciliation in a logarithmic number of rounds in order to reduce computational load. It builds upon the CPI approach, but unlike [MTZ03], does not attempt to find some $\bar{n}_\triangle \geq n_\triangle$ to craft a message containing sufficient information to compute the set difference. Instead the protocol fixes some $\bar{m} \in \mathbb{N}$ and attempts reconciliation using $\bar{m}$ as an approximation of $n_\triangle$. If reconciliation fails, X_0 and X_1 are partitioned and reconciliation of the partitions is attempted with the same value $\bar{m}$ to determine the message size. Because the size of the partitions decreases, so must eventually the size of the symmetric difference between the pairs of partitions, until it is less than $\bar{m}$ and reconciliation succeeds.

This approach eliminates the cubic scaling of the computation time in $\bar{n}_{\triangle}$. Similar to our auxiliary tree, the authors propose a tree structure of partitions where a parent node includes the intervals of its children. The reconciliation message for each of these intervals is precomputed and stored within the tree.

Given such a partition tree, the reconciliation procedure has the same asymptotic complexity bounds as ours. It has better constant terms: whereas our approach recurses whenever the two ranges that are being compared are unequal, i.e., whenever the size of the symmetric difference is greater than zero, their approach only recurses whenever the size of the symmetric difference is greater than $\bar{m}$.

The tree can be updated over insertions and deletions in time logarithmic in the height of the tree, but these updates are not balancing. In the worst case, maintaining this auxiliary data structure can thus take time linear in the size of the set per update. Furthermore, the tree is specific to a particular choice of evaluation and control points for the characteristic polynomial. If the tree is to be reused across different reconciliation sessions, these points have to be fixed in advance. This could allow an attacker to craft malicious sets for which a failed reconciliation is not detected. As a consequence, this approach can only be used with the producers of the sets are trusted.

Chapter 6

Conclusion

We have presented range-based set synchronization, an approach that can reconcile two sets in $\mathcal{O}(\log(n))$ communication rounds, transmitting $\mathcal{O}(\min(n_\triangle \cdot \log_b(n), n))$ bits, and performing computations in each communication round requiring $\mathcal{O}(\min(k \cdot \log(|S|), |S|))$ time and $\mathcal{O}(1)$ space. Other approaches in the literature improve on the communication complexity by removing the logarithmic factor, but at the cost of at least $\mathcal{O}(n)$ time and $\mathcal{O}(n_\triangle)$ space for the per-round computations. The range-based approach is easily adapted to mirror a set, none of the other set reconciliation literature can be modified in this way.

Low computational complexity and bounded space usage are required for robust and scalable implementations, transmitting logarithmically more bits is still more effective in practice than running out of memory and crashing. We thus believe that range based synchronization should be part of the toolkit for building distributed systems just like the more complex reconciliation protocols. As an extreme example, a microcontroller connected to network-attached storage could perform reconciliation of large sets even though it has restrictive computational capabilities.

For maximized efficiency in the face of constrained computational resources, hybrid approaches can be considered. The range-based approach transmits range item sets if the number of items inside a range falls below a threshold b. The simple exchange of these items can be considered as a minor optimization on the trivial reconciliation protocol, but any other set reconciliation protocol could also be used. Range-based fingerprint comparisons can thus be used to find ranges containing less than b items for both nodes, these ranges can then be reconciled using e.g. characteristic polynomial interpolation or Invertible Bloom Lookup Tables. Note however that if the space usage of the computations is to remain in $\mathcal{O}(1)$, these "leaf" reconciliations have to be performed sequentially, not in parallel.

6.1 Future Work

Beyond the scope of this thesis, there are some natural further questions to study. Section 2.4.1 lists some qualitative arguments for the validity of the approach even in the presence of a small number of adversarially created hash collisions. To give full confidence, a quantitative cryptographic analysis needs to answer questions such as the number of pairwise distinct collisions that can be adequately protected against through range boundary randomization, how well range boundary randomization can protect against k different subsets all hashing to the same digest, and whether

attacks against the collision resistance of the relevant monoids can be extended to produce such k-wise collisions.

We have considered the range-based approach for a two-party setting only, a natural generalization is that of multi-party set synchronization. This can of course be reduced to pairwise synchronization, but an interesting question is whether more efficient range-based protocols can be developed, similar to how [BM14] adapts characteristic polynomial interpolation and [MP18] adapts invertible bloom lookup tables.

While we have studied the protocol from a theoretical perspective, there are also some software design issues that arise in the context of actual implementations. Consider for example one of the original motivating examples, an unordered peer-to-peer pubsub mechanism. A node in such a network would ideally perform set reconciliation as a first step in a session, and then eagerly forward new items as they became available. Thus new items might arrive in the middle of another reconciliation step, the same problem occurs when doing multiple reconciliation steps in parallel. The new items need to be buffered until reconciliation has finished, and then they can be transmitted to the peer. One solution is to implement the sets as persistent data structures, but are there more efficient mutable solutions based on locking? Or can the buffering of concurrently arriving items possibly avoid it altogether? We have provided the theoretical basis for robust, scalable systems based on set replication, but a practical implementation raises interesting questions in its own right.

Bibliography

[BBJ+16] Michael A Bender, Jonathan W Berry, Rob Johnson, Thomas M Kroeger, Samuel McCauley, Cynthia A Phillips, Bertrand Simon, Shikha Singh, and David Zage. Anti-persistence on persistent storage: History-independent sparse tables and dictionaries. In *Proceedings of the 35th ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems*, pages 289–302, 2016.

[BCM02] John Byers, Jeffrey Considine, and Michael Mitzenmacher. Fast approximate reconciliation of set differences. In *BU Computer Science TR*. Citeseer, 2002.

[Ben75] Jon Louis Bentley. Multidimensional binary search trees used for associative searching. *Communications of the ACM*, 18(9):509–517, 1975.

[BLL02] Ahto Buldas, Peeter Laud, and Helger Lipmaa. Eliminating counterevidence with applications to accountable certificate management 1. *Journal of Computer Security*, 10(3):273–296, 2002.

[Blo70] Burton H Bloom. Space/time trade-offs in hash coding with allowable errors. *Communications of the ACM*, 13(7):422–426, 1970.

[BM97] Mihir Bellare and Daniele Micciancio. A new paradigm for collision-free hashing: Incrementality at reduced cost. In *International Conference on the Theory and Applications of Cryptographic Techniques*, pages 163–192. Springer, 1997.

[BM02] Rudolf Bayer and Edward McCreight. Organization and maintenance of large ordered indexes. In *Software pioneers*, pages 245–262. Springer, 2002.

[BM14] Anudhyan Boral and Michael Mitzenmacher. Multi-party set reconciliation using characteristic polynomials. In *2014 52nd Annual Allerton Conference on Communication, Control, and Computing (Allerton)*, pages 1182–1187. IEEE, 2014.

[BSV17] Lisa Bromberg, Vladimir Shpilrain, and Alina Vdovina. Navigating in the cayley graph of $sl_2(\mathbb{F}_p)$ and applications to hashing. In *Semigroup Forum*, volume 94, pages 314–324. Springer, 2017.

[CDVD+03] Dwaine Clarke, Srinivas Devadas, Marten Van Dijk, Blaise Gassend, and G Edward Suh. Incremental multiset hash functions and their application to memory integrity checking. In *International conference*

on the theory and application of cryptology and information security, pages 188–207. Springer, 2003.

[CEG+99] Yuan Chen, Jan Edler, Andrew Goldberg, Allan Gottlieb, Sumeet Sobti, and Peter Yianilos. A prototype implementation of archival intermemory. In *Proceedings of the fourth ACM conference on Digital libraries*, pages 28–37, 1999.

[CL02] Jan Camenisch and Anna Lysyanskaya. Dynamic accumulators and application to efficient revocation of anonymous credentials. In *Annual International Cryptology Conference*, pages 61–76. Springer, 2002.

[CNQ09] Julien Cathalo, David Naccache, and Jean-Jacques Quisquater. Comparing with rsa. In *IMA International Conference on Cryptography and Coding*, pages 326–335. Springer, 2009.

[CS14] Scott Chacon and Ben Straub. *Pro git*. Springer Nature, 2014.

[Dev90] Luc Devroye. On the height of random m-ary search trees. *Random Structures & Algorithms*, 1(2):191–203, 1990.

[Dwo15] Morris J Dworkin. Sha-3 standard: Permutation-based hash and extendable-output functions. 2015.

[EGUV11] David Eppstein, Michael T Goodrich, Frank Uyeda, and George Varghese. What's the difference? efficient set reconciliation without prior context. *ACM SIGCOMM Computer Communication Review*, 41(4):218–229, 2011.

[Fen94] Peter M Fenwick. A new data structure for cumulative frequency tables. *Software: Practice and experience*, 24(3):327–336, 1994.

[GIMS11] Markus Grassl, Ivana Ilić, Spyros Magliveras, and Rainer Steinwandt. Cryptanalysis of the tillich–zúmor hash function. *Journal of cryptology*, 24(1):148–156, 2011.

[GL12] Deke Guo and Mo Li. Set reconciliation via counting bloom filters. *IEEE Transactions on Knowledge and Data Engineering*, 25(10):2367–2380, 2012.

[GM11] Michael T Goodrich and Michael Mitzenmacher. Invertible bloom lookup tables. In *2011 49th Annual Allerton Conference on Communication, Control, and Computing (Allerton)*, pages 792–799. IEEE, 2011.

[Gol09] Daniel Golovin. B-treaps: A uniquely represented alternative to b-trees. In *International Colloquium on Automata, Languages, and Programming*, pages 487–499. Springer, 2009.

[GT00] Michael T Goodrich and Roberto Tamassia. Efficient authenticated dictionaries with skip lists and commutative hashing. Technical report, Technical Report, Johns Hopkins Information Security Institute, 2000.

[HHM+05] Jason D Hartline, Edwin S Hong, Alexander E Mohr, William R Pentney, and Emily C Rocke. Characterizing history independent data structures. *Algorithmica*, 42(1):57–74, 2005.

[Jon03] Simon Peyton Jones. *Haskell 98 language and libraries: the revised report*. Cambridge University Press, 2003.

[LGR+19] Lailong Luo, Deke Guo, Ori Rottenstreich, Richard TB Ma, and Xueshan Luo. Set reconciliation with cuckoo filters. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management*, pages 2465–2468, 2019.

[LKMW19] Kevin Lewi, Wonho Kim, Ilya Maykov, and Stephen Weis. Securing update propagation with homomorphic hashing. *IACR Cryptol. ePrint Arch.*, 2019:227, 2019.

[Lyu05] Vadim Lyubashevsky. The parity problem in the presence of noise, decoding random linear codes, and the subset sum problem. In *Approximation, randomization and combinatorial optimization. Algorithms and techniques*, pages 378–389. Springer, 2005.

[MEL+21] Krzysztof Maziarz, Tom Ellis, Alan Lawrence, Andrew Fitzgibbon, and Simon Peyton Jones. Hashing modulo alpha-equivalence. *arXiv preprint arXiv:2105.02856*, 2021.

[Mer89] Ralph C Merkle. A certified digital signature. In *Conference on the Theory and Application of Cryptology*, pages 218–238. Springer, 1989.

[MGS15] Hristina Mihajloska, Danilo Gligoroski, and Simona Samardjiska. Reviving the idea of incremental cryptography for the zettabyte era use case: Incremental hash functions based on sha-3. In *International Workshop on Open Problems in Network Security*, pages 97–111. Springer, 2015.

[MND+04] Charles Martel, Glen Nuckolls, Premkumar Devanbu, Michael Gertz, April Kwong, and Stuart G Stubblebine. A general model for authenticated data structures. *Algorithmica*, 39(1):21–41, 2004.

[MP18] Michael Mitzenmacher and Rasmus Pagh. Simple multi-party set reconciliation. *Distributed Computing*, 31(6):441–453, 2018.

[MSTA17] Jeremy Maitin-Shepard, Mehdi Tibouchi, and Diego F Aranha. Elliptic curve multiset hash. *The Computer Journal*, 60(4):476–490, 2017.

[MT02] Yaron Minsky and Ari Trachtenberg. Practical set reconciliation. In *40th Annual Allerton Conference on Communication, Control, and Computing*, volume 248. Citeseer, 2002.

[MT16] Ciaran Mullan and Boaz Tsaban. sl_2 homomorphic hash functions: worst case to average case reduction and short collision search. *Designs, Codes and Cryptography*, 81(1):83–107, 2016.

[MTZ03] Yaron Minsky, Ari Trachtenberg, and Richard Zippel. Set reconciliation with nearly optimal communication complexity. *IEEE Transactions on Information Theory*, 49(9):2213–2218, 2003.

[MVOV18] Alfred J Menezes, Paul C Van Oorschot, and Scott A Vanstone. *Handbook of applied cryptography.* CRC press, 2018.

[NN00] Moni Naor and Kobbi Nissim. Certificate revocation and certificate update. *IEEE Journal on selected areas in communications*, 18(4):561–570, 2000.

[OAL+19] A Pinar Ozisik, Gavin Andresen, Brian N Levine, Darren Tapp, George Bissias, and Sunny Katkuri. Graphene: efficient interactive set reconciliation applied to blockchain propagation. In *Proceedings of the ACM Special Interest Group on Data Communication*, pages 303–317. 2019.

[Pet09] Christophe Petit. *On graph-based cryptographic hash functions.* PhD thesis, Catholic University of Louvain, Louvain-la-Neuve, Belgium, 2009.

[PQ10] Christophe Petit and Jean-Jacques Quisquater. Preimages for the tillich-zémor hash function. In *International Workshop on Selected Areas in Cryptography*, pages 282–301. Springer, 2010.

[PQ+11] Christophe Petit, Jean-Jacques Quisquater, et al. Rubik's for cryptographers. *IACR Cryptol. ePrint Arch.*, 2011:638, 2011.

[PSTY13] Charalampos Papamanthou, Elaine Shi, Roberto Tamassia, and Ke Yi. Streaming authenticated data structures. In *Annual International Conference on the Theory and Applications of Cryptographic Techniques*, pages 353–370. Springer, 2013.

[PT89] William Pugh and Tim Teitelbaum. Incremental computation via function caching. In *Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 315–328, 1989.

[PTT11] Charalampos Papamanthou, Roberto Tamassia, and Nikos Triandopoulos. Optimal verification of operations on dynamic sets. In *Annual Cryptology Conference*, pages 91–110. Springer, 2011.

[PTT16] Charalampos Papamanthou, Roberto Tamassia, and Nikos Triandopoulos. Authenticated hash tables based on cryptographic accumulators. *Algorithmica*, 74(2):664–712, 2016.

[Pug90] William Pugh. Skip lists: a probabilistic alternative to balanced trees. *Communications of the ACM*, 33(6):668–676, 1990.

[Pug98] William Pugh. A skip list cookbook. Technical report, 1998.

[SA96] Raimund Seidel and Cecilia R Aragon. Randomized search trees. *Algorithmica*, 16(4):464–497, 1996.

[SMBA10] Paul T Stanton, Benjamin McKeown, Randal Burns, and Giuseppe Ateniese. Fastad: an authenticated directory for billions of objects. *ACM SIGOPS Operating Systems Review*, 44(1):45–49, 2010.

[Sny77] L. Snyder. On uniquely represented data strauctures. In *2013 IEEE 54th Annual Symposium on Foundations of Computer Science*, pages 142–146, Los Alamitos, CA, USA, oct 1977. IEEE Computer Society. URL: `https://doi.ieeecomputersociety.org/10.1109/SFCS.1977.22`, `doi:10.1109/SFCS.1977.22`.

[Sos16] Bianca Sosnovski. Cayley graphs of semigroups and applications to hashing. 2016.

[ST94] Rajamani Sundar and Robert E Tarjan. Unique binary-search-tree representations and equality testing of sets and sequences. *SIAM Journal on Computing*, 23(1):24–44, 1994.

[SYW+17] Wentao Shang, Yingdi Yu, Lijing Wang, Alexander Afanasyev, and Lixia Zhang. A survey of distributed dataset synchronization in named data networking. *NDN, Technical Report NDN-0053*, 2017.

[TLMT19] Dominic Tarr, Erick Lavoie, Aljoscha Meyer, and Christian Tschudin. Secure scuttlebutt: An identity-centric protocol for subjective and decentralized applications. In *Proceedings of the 6th ACM Conference on Information-Centric Networking*, pages 1–11, 2019.

[TM+96] Andrew Tridgell, Paul Mackerras, et al. The rsync algorithm. 1996.

[TZ94] Jean-Pierre Tillich and Gilles Zémor. Hashing with sl 2. In *Annual International Cryptology Conference*, pages 40–49. Springer, 1994.

[TZX+11] Xiaomei Tian, Dafang Zhang, Kun Xie, Can Hu, Mengfan Wang, and Jinguo Deng. Exact set reconciliation based on bloom filters. In *Proceedings of 2011 International Conference on Computer Science and Network Technology*, volume 3, pages 2001–2009. IEEE, 2011.

[Wag02] David Wagner. A generalized birthday problem. In *Annual International Cryptology Conference*, pages 288–304. Springer, 2002.

[WC81] Mark N Wegman and J Lawrence Carter. New hash functions and their use in authentication and set equality. *Journal of computer and system sciences*, 22(3):265–279, 1981.

[Zém91] Gilles Zémor. Hash functions and graphs with large girths. In *Workshop on the Theory and Application of of Cryptographic Techniques*, pages 508–511. Springer, 1991.